arXiv:2610.40276v1 [cs.CV] 29 Oct 2026

TaskFrame: Connecting Task-Level Planning with Frame-Level Control for Game Agents

Guangzhi Wang[1*], **Junfu Pu**[1*†], **Yiran Wang**[1*], **Xingyilang Yin**[2*], **Honghao Chen**[1], **Cheng Cheng**[1], **Lingen Li**[3], **Yinan Zhou**[4], **Jiarui Wu**[3], **Yinuo Jiang**[5], **Hongbo Jin**[6], **Xiaodong Cun**[1,7‡], **Ying Shan**[1]

[1]Tencent ARC Lab, [2]University of Macau, [3]The Chinese University of Hong Kong, [4]City University of Hong Kong, [5]Huazhong University of Science and Technology, [6]Peking University, [7]Great Bay University
[*]Equal contribution., [†]Project lead., [‡]Corresponding author.

Building general game agents requires both high-level task reasoning and low-latency action control in rapidly changing visual environments. These capabilities favor models with substantially different computational profiles, making them difficult to realize within a single model. In this work, we present TaskFrame, a hierarchical planner–controller framework that connects task-level planning with frame-level keyboard-and-mouse control. A vision-language task planner decomposes a high-level goal into an ordered sequence of atomic instructions and is optimized through supervised learning and dependency-aware GRPO. The hidden states of special goal tokens that terminate each atomic instruction provide continuous goal embeddings that are cached and reused across control steps. A causal Transformer controller combines the embedding of the active goal with recent observations to predict discrete keyboard and mouse-button actions and continuous camera motion at every game step. We train the system in three stages: planner training, controller pretraining, and joint alignment. Experiments in Minecraft show that TaskFrame is competitive with general-purpose vision-language models on instruction decomposition, and that in closed-loop evaluation the agent achieves a higher combat success rate than the evaluated baselines.

Date: October 30, 2026
Project page: https://task-frame.github.io/

1 Introduction

Games have long been a demanding testbed for artificial general intelligence (AGI), as successful gameplay requires the integration of perception, reasoning, planning, and decision making into goal-directed behavior. Existing efforts have reached or surpassed expert human performance in board games, Atari, multiplayer poker, and complex real-time strategy games (Silver et al., 2016, 2017; Schrittwieser et al., 2020; OpenAI et al., 2019; Badia et al., 2020; Brown and Sandholm, 2019; Vinyals et al., 2019). More recently, research has shifted toward *embodied game agents* that follow language instructions and act directly from pixels in open-ended virtual worlds (SIMA Team, 2025; Magne et al., 2026; Yue et al., 2026; Li et al., 2025a; Hafner et al., 2025). Unlike agents specialized for a fixed environment or objective, such systems provide a natural setting for studying whether an agent can understand diverse goals, adapt its behavior to changing visual situations, and execute complex tasks through a general action interface.

Building such an embodied agent remains difficult because task-level reasoning and frame-level control operate at different frequencies. Understanding a scene, grounding an instruction, and determining the correct sequence of subgoals require a large, high-capacity model. In contrast, executing those subgoals requires a lightweight controller that can continuously incorporate the latest frame and emit actions with low latency. This challenge is especially evident in visually dynamic games, where abrupt camera motion and sudden events must be handled through a hybrid action space combining discrete keyboard presses with continuous mouse movement. Running a large reasoning model at every control step is prohibitively expensive, whereas a small reactive policy alone lacks the semantic capacity to reliably interpret and organize complex instructions and visual history.

arXiv:2610.40276v1 [cs.CV] 29 Oct 2026

TaskFrame：连接游戏智能体的任务级规划与帧级控制

Guangzhi Wang[1*], **Junfu Pu**[1*†], **Yiran Wang**[1*], **Xingyilang Yin**[2*], **Honghao Chen**[1], **Cheng Cheng**[1], **Lingen Li**[3], **YinanZhou**[4], **Jiarui Wu**[3], **Yinuo Jiang**[5], **Hongbo Jin**[6], **XiaodongCun**[1, 7‡], **Ying Shan**[1]

[1]腾讯ARC实验室， [2]澳门大学， [3]香港中文大学， [4]香港城市大学， [5]华中科技大学， [6]北京大学， [7]大湾区大学
[*]同等贡献。， [†]项目负责人。， [‡]通讯作者。

构建通用游戏智能体，需要在快速变化的视觉环境中同时具备高层任务推理与低延迟动作控制能力。这两种能力所适用的模型在计算特性上存在显著差异，因此难以在单一模型中兼顾。在本研究中，我们提出了TaskFrame，这是一种连接任务级规划与帧级键鼠控制的分层规划器–控制器框架。视觉语言任务规划器将高层目标分解为有序的原子指令序列，并通过监督学习和依赖感知GRPO进行优化。每条原子指令末尾的特殊目标词元的隐藏状态提供连续目标嵌入，这些嵌入会被缓存，并在各控制步中重复使用。因果Transformer控制器将当前目标的嵌入与近期观测相结合，在每个游戏步预测离散的键盘与鼠标按键动作以及连续相机运动。我们分三个阶段训练该系统：规划器训练、控制器预训练和联合对齐。在我的世界中的实验表明，TaskFrame在指令分解方面可与通用视觉语言模型相媲美；在闭环评估中，该智能体的战斗成功率高于所评估的基线。

日期：2026年10月30日**项目主页：**
https://task-frame.github.io/

1 引言

长期以来，游戏一直是通用人工智能（AGI）极具挑战性的试验平台，因为成功的游戏过程需要将感知、推理、规划和决策整合为目标导向行为。现有研究已在棋盘游戏、雅达利、多人扑克和复杂的即时战略游戏中达到或超越人类专家水平（Silver 等人，2016年，2017年；Schrittwieser 等人，2020年；OpenAI 等人，2019年；Badia 等人，2020年；Brown 和 Sandholm，2019年；Vinyals 等人，2019年）。近年来，研究重心已转向具身游戏智能体，这类智能体遵循语言指令，并在开放式虚拟世界中直接依据像素采取行动（SIMA团队，2025年；Magne 等人，2026年；Yue等人，2026年；Li 等人，2025a；Hafner 等人，2025年）。与专门针对固定环境或目标的智能体不同，此类系统提供了一个自然的研究场景，用于探究智能体能否理解多样化目标、根据不断变化的视觉情境调整自身行为，并通过通用动作接口执行复杂任务。

构建这样的具身智能体仍然十分困难，因为任务级推理与帧级控制以不同的频率运行。理解场景、实现指令落地以及确定正确的子目标序列，都需要规模大、能力强的模型。相比之下，执行这些子目标则需要轻量级控制器，能够持续纳入最新帧并以低延迟输出动作。这一挑战在视觉动态游戏中尤为突出：智能体必须通过结合离散键盘按键操作与连续鼠标移动的混合动作空间，来应对突发的相机运动和突如其来的事件。在每个控制步都运行大型推理模型的成本高得难以承受，而仅靠小型反应式策略又缺乏足够的语义能力，无法可靠地理解和组织复杂指令及视觉历史。

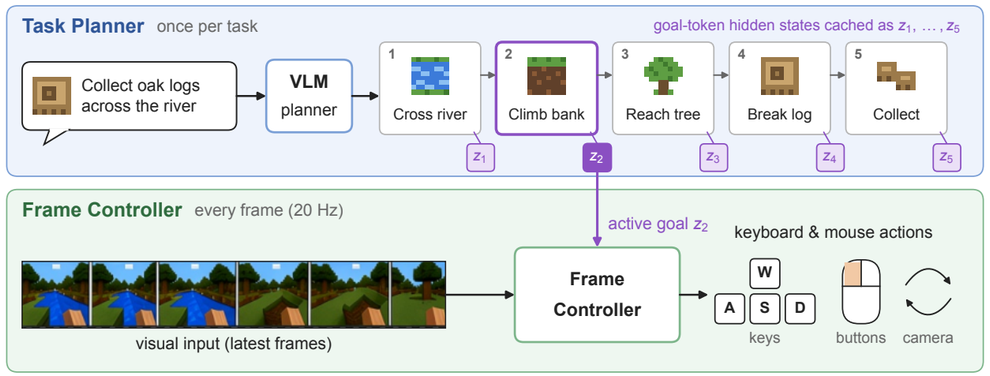

Figure 1 Overview of TaskFrame. Once per task, the task planner decomposes a high-level instruction into ordered atomic goals, each terminated by a special goal token whose hidden state is cached as a goal embedding. At every frame, the frame controller takes the embedding of the active goal and the latest visual input and emits native keyboard-and-mouse actions.

Existing work addresses parts of this tension through several strategies. Hierarchical agents use a language model to select skills or produce high-level plans for a goal-conditioned controller (Ahn et al., 2022; Wang et al., 2023c; Yuan et al., 2023; Li et al., 2025b). In physical robotics, dual-system Vision-Language-Action (VLA) models separate a slower semantic module from a faster visuomotor policy (NVIDIA et al., 2025; Figure AI, 2025; Zhang et al., 2024; Bu et al., 2024; Shi et al., 2025), while action chunking amortizes inference by predicting multiple actions at once, with recent methods improving parallel decoding or asynchronous execution (Kim et al., 2025; Black et al., 2025). A practical game agent must connect these mechanisms through an interface that supports both structured planning and responsive control. The high-level model must translate an open-vocabulary goal into a complete and correctly ordered sequence of atomic goals, and the controller must keep incorporating the latest visual observation at every action step.

We present **TaskFrame**, a hierarchical game agent that connects task-level planning with frame-level control through planner-produced goal representations (Figure 1). Its high-capacity *task planner* is trained to perform visually grounded instruction decomposition, with dedicated data and a dependency-aware reinforcement learning objective that addresses both step-count and ordering correctness. Each atomic goal it emits is terminated by a special goal token, and the hidden states at these tokens are cached and passed to a lightweight *frame controller* as continuous goal embeddings. The controller combines the embedding of the active goal with recent frames to emit discrete keyboard and mouse-button actions and continuous camera motion at every game step, so the large planner stays outside the frame-level control loop. We train the system in three stages: the planner is trained for decomposition, the controller is trained to produce actions from each frame, and the two modules are finally aligned jointly.

We train the controller on human gameplay with native keyboard-and-mouse actions recorded directly from players, and the planner on the hierarchical annotations of GameHorizon Suite (Wang et al., 2026), which provide atomic instructions, compositional goals, and dependency relations. Our 9B planner outperforms GPT-4o and Qwen3.5-397B-A17B on instruction decomposition. In closed-loop Minecraft evaluation, TaskFrame reaches a 51% average success rate on combat tasks, the highest among the evaluated models, and ablations examine each component. To summarize, our contributions are as follows:

- We introduce TaskFrame, a hierarchical game agent in which a vision-language planner decomposes goals once per task and conditions a frame-level keyboard-and-mouse controller, which acts at every game step, through the hidden states of special goal tokens.
- We propose a dependency-aware GRPO objective for instruction decomposition that rewards necessary precedence relations separately from conventional ordering, derived from pairwise dependency

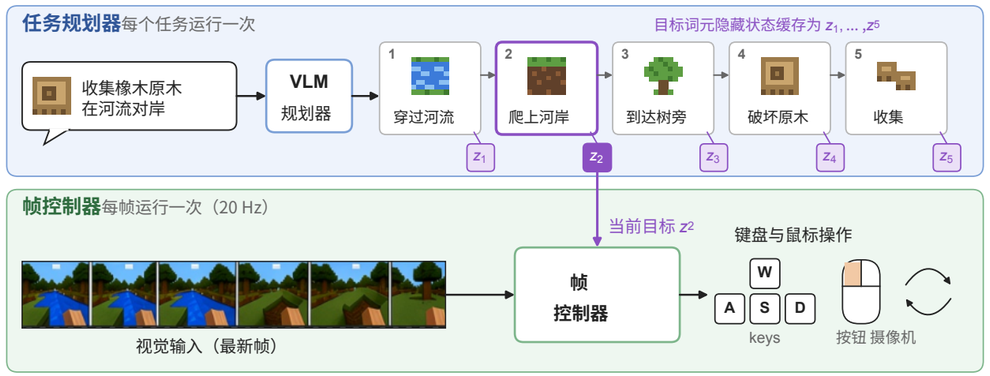

图 1 TaskFrame 概览。任务规划器针对每个任务执行一次规划，将高层指令分解为有序的原子目标，每个原子目标均以一个特殊的目标词元结束，该词元的隐藏状态被缓存为目标嵌入。在每一帧，帧控制器接收当前目标的嵌入和最新的视觉输入，并输出原生键鼠动作。

现有研究通过多种策略部分缓解了这一矛盾。分层智能体使用语言模型来选择技能，或为目标条件控制器生成高层规划（Ahn 等人，2022年；Wang 等人，2023c；Yuan 等人，2023年；Li 等人，2025b）。在实体机器人技术领域，双系统视觉-语言-动作（VLA）模型将较慢的语义模块与较快的视觉运动策略分离（英伟达等人，2025年；Figure AI，2025年；Zhang 等人，2024年；Bu 等人，2024年；Shi 等人，2025年），而动作分块通过一次预测多个动作来分摊推理成本，近期方法则进一步改进了并行解码或异步执行（Kim 等人，2025年；Black 等人，2025年）。实用的游戏智能体必须通过同时支持结构化规划和响应式控制的接口，将这些机制连接起来。高层模型必须将开放词汇目标转化为完整且顺序正确的原子目标序列，而控制器则必须在每个动作步持续纳入最新的视觉观测。

我们提出了 **TaskFrame**，这是一种分层游戏智能体，通过规划器生成的目标表征，将任务级规划与帧级控制连接起来（图 1）。其高容量任务规划器利用专门的数据和依赖感知强化学习目标进行训练，以执行视觉锚定的指令分解，同时兼顾步骤数量与顺序的正确性。它输出的每个原子目标均以一个特殊的目标词元结束，这些词元处的隐藏状态会被缓存，并作为连续目标嵌入传递给轻量级的帧控制器。控制器将当前目标的嵌入与最近的帧相结合，在每个游戏步输出离散的键盘与鼠标按键动作以及连续相机运动，从而使大型规划器不必参与帧级控制回路。我们分三个阶段训练该系统：首先训练规划器进行分解，然后训练控制器根据每一帧生成动作，最后对两个模块进行联合对齐。

我们使用直接记录玩家原生键鼠动作的人类游戏操作数据来训练控制器，并使用 GameHorizon Suite（Wang 等人，2026年）的分层标注来训练规划器，这些标注提供了原子指令、组合式目标和依赖关系。我们的90亿参数规划器在指令分解方面优于 GPT-4o 和 Qwen3.5-397B-A17B。在Minecraft闭环评估中，TaskFrame 在战斗任务上达到了 51% 的平均成功率，在所有参评模型中最高；我们还通过消融实验考察了各个组件。总之，我们的贡献如下：

- 我们提出了 TaskFrame，这是一种分层游戏智能体，其中的视觉语言规划器针对每项任务进行一次目标分解，并通过特殊目标词元的隐藏状态，为在每个游戏步执行操作的帧级键鼠控制器提供条件信息。
- 我们提出了一种用于指令分解的依赖感知的GRPO目标函数，将必要的先后依赖关系与惯常顺序分开给予奖励，其依据来自成对依赖关系

annotations.

- We propose a three-stage training recipe that pretrains the controller with a generic goal encoder and then aligns it with the planner through LoRA adapters, and we equip the controller with a hybrid discrete and flow-matching action head. We evaluate the planner on two decomposition benchmarks and the full agent in closed-loop Minecraft control.

2 Related Work

Game agents. MineRL and VPT established large-scale imitation from human Minecraft play (Guss et al., 2019; Baker et al., 2022), and MineDojo added an open-ended environment with internet-scale knowledge (Fan et al., 2022). STEVE-1, GROOT, ROCKET-1, and MineDreamer condition low-level policies on text, reference videos, visual-temporal prompts, or imagined goals (Lifshitz et al., 2023; Cai et al., 2024b, 2025; Zhou et al., 2024), while JARVIS-VLA and OmniJARVIS post-train VLMs to emit keyboard-and-mouse actions directly (Li et al., 2025a; Wang et al., 2024), following robotic VLAs (Brohan et al., 2023; Kim et al., 2024). Recent systems generalize across tasks and games (Wang et al., 2025; SIMA Team, 2025; Tan et al., 2025, 2024) and control games through native keyboard-and-mouse interfaces at frame rate (Yue et al., 2026; Magne et al., 2026). TaskFrame instead connects explicit, trained task decomposition to frame-level control and keeps the large VLM outside the per-frame loop.

Hierarchical task planning. Language-model planners ground decisions in skill affordances (Ahn et al., 2022), use environment feedback (Huang et al., 2022), or write code over control primitives (Liang et al., 2023). In Minecraft, Plan4MC, DEPS, Voyager, JARVIS-1, and Optimus-1/2 plan over learned skills or dispatch multimodal plans to goal-conditioned controllers (Yuan et al., 2023; Wang et al., 2023b,a,c; Li et al., 2024, 2025b). These planners are prompted or supervised to produce plans but are not optimized for whether their steps respect task dependencies. Reinforcement learning with verifiable rewards, popularized by GRPO and its variants (Shao et al., 2024; DeepSeek-AI, 2025; Liu et al., 2025), offers a way to do so: we use the must-precede annotations of GameHorizon Suite (Wang et al., 2026) to define a reward that separates necessary precedence from conventional ordering.

Dual-system control. Dual-system policies pair a slow semantic module with a fast visuomotor policy (Kahneman, 2011). Hi Robot and RT-H pass language between the two (Shi et al., 2025; Belkhale et al., 2024), and $\pi_{0.5}$ predicts a textual subtask before decoding actions (Physical Intelligence et al., 2025); HiRT, RoboDual, GR00T N1, and Helix pass latent features from a slowly updated VLM (Zhang et al., 2024; Bu et al., 2024; NVIDIA et al., 2025; Figure AI, 2025), LCB trains a dedicated latent token as the bridge (Shentu et al., 2024), and Acting While Understanding updates semantic conditions asynchronously (Yan et al., 2026). Action chunking, better action tokenization, and overlapping generation with execution further improve control efficiency (Zhao et al., 2023; Chi et al., 2025; Black et al., 2024; Pertsch et al., 2025; Kim et al., 2025; Black et al., 2025; Shukor et al., 2025). Like HiRT and LCB, TaskFrame passes latent goals from a VLM to a fast policy, but its planner is first optimized for ordered decomposition, and its latent goals are the hidden states of goal tokens emitted within that decomposition, so one planner pass yields both an interpretable plan and one embedding per step.

3 Method

TaskFrame connects task-level planning with frame-level game control (Figure 2). A vision-language *task planner* $\mathcal{B}$ maps an observation and a high-level instruction to an ordered sequence of atomic goals, each terminated by a special goal token τ whose hidden state serves as a continuous representation of that goal. A causal *frame controller* $\mathcal{P}$ consumes the representation of the active goal together with its recent frame and action history and predicts native keyboard-and-mouse actions at every game step.

Training proceeds in three stages. First, we train the planner for visually grounded instruction decomposition with supervised fine-tuning (SFT) and dependency-aware GRPO (Section 3.1). Second, we pretrain the controller on atomic-segment gameplay, using a frozen, generic vision-language model as its goal encoder

标注。

- 我们提出了一套三阶段训练方案，先使用通用目标编码器预训练控制器，再通过LoRA适配器使其与规划器对齐，并为控制器配备离散与流匹配混合动作头。我们在两个分解基准上评估规划器，并在Minecraft闭环控制中评估完整智能体。

2 相关工作

游戏智能体。MineRL和VPT奠定了基于人类《我的世界》游戏数据进行大规模模仿学习的基础（Guss 等人，2019年；Baker 等人，2022年），MineDojo则引入了融合互联网规模知识的开放式环境（Fan 等人，2022年）。STEVE-1、GROOT、ROCKET-1和MineDreamer以文本、参考视频、视觉时序提示或想象目标作为低层策略的条件（Lifshitz 等人，2023年；Cai 等人，2024b，2025年；Zhou 等人，2024年），而JARVIS-VLA和OmniJARVIS沿用机器人视觉语言动作模型的思路（Brohan 等人，2023年；Kim 等人，2024年），对视觉语言模型进行后训练，使其直接输出键鼠动作（Li 等人，2025a；Wang 等人，2024年）。近期系统能够跨任务和游戏泛化（Wang 等人，2025年；SIMA团队，2025年；Tan 等人，2025年，2024年），并通过原生键鼠接口以帧率控制游戏（Yue等人，2026年；Magne 等人，2026年）。相比之下，TaskFrame将经过训练的显式任务分解与帧级控制连接起来，并将大型视觉语言模型置于逐帧循环之外。

分层任务规划。语言模型规划器基于技能可供性做出决策（Ahn 等人，2022年），利用环境反馈（Huang 等人，2022年），或基于控制原语编写代码（Liang 等人，2023年）。在《我的世界》中，Plan4MC、DEPS、Voyager、JARVIS-1 和 Optimus-1/2 基于习得技能进行规划，或将多模态计划分派给目标条件控制器（Yuan 等人，2023年；Wang 等人，2023b,a,c；Li 等人，2024年，2025b）。这些规划器通过提示词或监督来生成计划，但并未针对计划步骤是否遵循任务依赖关系进行优化。由组相对策略优化及其变体（Shao 等，2024年；DeepSeek-AI，2025年；Liu等人，2025年）推广的可验证奖励强化学习为此提供了一种途径：我们利用 GameHorizon Suite（Wang 等人，2026年）的必须先行关系标注来定义奖励，以区分必要先后关系与惯常顺序。

双系统控制。双系统策略将慢速语义模块与快速视觉运动策略配对（Kahneman，2011年）。Hi Robot 和 RT-H 在两者之间传递语言（Shi 等，2025年；Belkhale 等，2024年），而$\pi_{0.5}$ 在解码动作之前预测文本子任务（Physical Intelligence 等，2025年）；HiRT、RoboDual、GR00T N1 和 Helix 传递来自缓慢更新的视觉语言模型的潜在特征（Zhang 等，2024年；Bu 等，2024年；英伟达等，2025年；Figure AI，2025年），LCB 训练一个专用的潜在词元作为桥梁（Shentu 等，2024年），Acting While Understanding 则异步更新语义条件（Yan 等，2026年）。动作分块、更优的动作词元化以及生成与执行的重叠进一步提高了控制效率（Zhao 等，2023年；Chi 等，2025年；Black 等，2024年；Pertsch 等，2025年；Kim 等，2025年；Black 等，2025年；Shukor 等，2025年）。与 HiRT 和 LCB 一样，TaskFrame 将潜在目标从视觉语言模型传递给快速策略，但其规划器首先针对有序分解进行优化，其潜在目标则是该分解过程中输出的目标词元的隐藏状态，因此，规划器的一次前向计算即可同时生成一个可解释计划和每个步骤对应的一个嵌入。

3 方法

TaskFrame 将任务级规划与帧级游戏控制相连接（图 2）。视觉语言任务规划器 $\mathcal{B}$ 将观测和高层指令映射为有序的原子目标序列，每个原子目标均以一个特殊的目标词元 τ 结束，该词元的隐藏状态作为该目标的连续表示。因果帧控制器 $\mathcal{P}$ 接收当前目标的表示及其近期的帧和动作历史，并在每个游戏步预测原生键鼠动作。

训练分为三个阶段。首先，我们通过监督微调（SFT）和依赖感知GRPO训练规划器，使其具备视觉锚定的指令分解能力（第3.1节）。其次，我们在原子片段游戏过程上预训练控制器，使用冻结的通用视觉语言模型作为其目标编码器

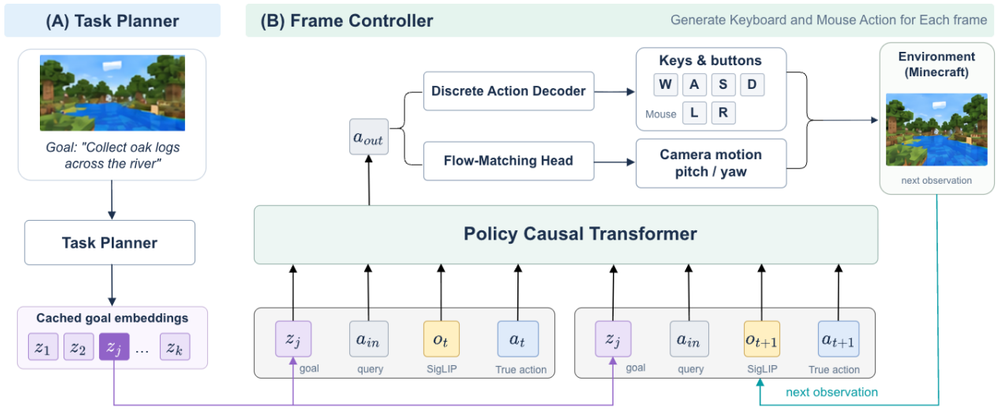

Figure 2 Overview of TaskFrame, a hierarchical planner–controller agent that separates task-level planning from frame-level control. **(A) Task Planner.** A vision–language model decomposes a high-level goal (e.g., *"collect oak logs across the river"*) into an ordered sequence of atomic goals, each terminated by a special goal token. The hidden states at these tokens are cached as goal embeddings $z_1, \ldots, z_k$. **(B) Frame Controller.** A per-frame autoregressive policy. Each frame consumes the active goal embedding z_j, an action-query token a_{in}, a SigLIP image token o_t, and the action a_t, which only later frames attend to. A causal Transformer aggregates them into the action-output token a_{out}, which is decoded into discrete keyboard and mouse-button presses and continuous camera motion. The planner runs once per high-level goal, while the controller acts at every game step.

(Section 3.2). Third, we replace this encoder with the trained planner and jointly adapt the controller and LoRA adapters in the planner with the action objective (Section 3.3). At control step t, the controller observes a frame $o_t \in \mathbb{R}^{H\times W\times 3}$ and predicts

$$a_t = (a_t^{\text{key}}, a_t^{\text{btn}}, a_t^{\text{mouse}}), \tag{1}$$

where the first two components are discrete actions and the last is the continuous camera motion produced by the mouse, represented as yaw and pitch changes $a_t^{\text{mouse}} = (\Delta\psi_t, \Delta\theta_t) \in \mathbb{R}^2$ in degrees.

3.1 Task Planner

Visually grounded decomposition. The task planner $\mathcal{B}$ converts a visual observation and a high-level instruction into an ordered sequence of atomic goals. Given an observation o and a high-level instruction g, it generates

$$\hat{\mathcal{S}} = (\hat{s}_1, \ldots, \hat{s}_{\hat{k}}) \sim \pi_\theta(\cdot \mid o, g), \tag{2}$$

where each $\hat{s}_j$ is an atomic instruction. We train the planner on the hierarchical gameplay annotations of GameHorizon Suite (Wang et al., 2026), which associate high-level goals with temporally aligned atomic instructions and annotate precedence dependencies between them. During training, the target number of goals k is supplied through a fixed reasoning prefix, yielding $\pi_\theta(\cdot \mid o, g, k)$; at inference it is not supplied.

Starting from Qwen3.5-9B, we freeze the vision encoder and train the language model and multimodal connector on the reference output sequence $\mathbf{y}$, whose goals are each terminated by a goal token (Equation (11)):

$$\mathcal{L}_{\text{SFT}} = -\sum_{t=1}^{|\mathbf{y}|} \log \pi_\theta(y_t \mid o, g, k, y_{<t}). \tag{3}$$

This stage teaches the planner to describe atomic goals in a consistent decomposition format and provides the initial policy for subsequent refinement. Data statistics, hyperparameters, and an analysis of why naive SFT overfits to annotation phrasing are given in Section A.

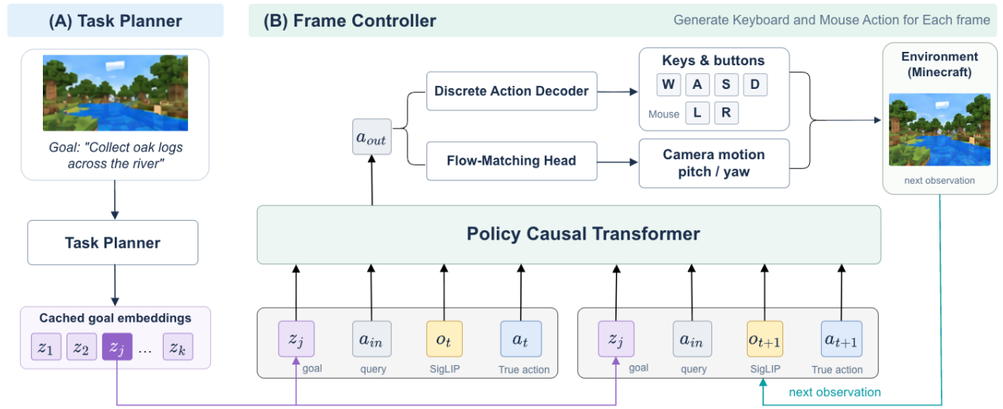

图2 TaskFrame概览，这是一种将任务级规划与帧级控制分离的分层规划器–控制器智能体。**（A）任务规划器**。视觉语言模型将高层目标（例如"到河对岸收集橡木原木"）分解为有序的原子目标序列，每个原子目标均以一个特殊的目标词元结束。这些词元处的隐藏状态被缓存为目标嵌入 z_1，……，z_k。**（B）帧控制器**。一种逐帧自回归策略。每帧接收当前目标嵌入 z_j、动作查询词元 a_{in}、SigLIP图像词元 o_t以及动作a_t，其中动作仅供后续帧关注。因果Transformer将它们聚合为动作输出词元 a_{out}，再将其解码为离散的键盘及鼠标按键操作和连续相机运动。规划器针对每个高层目标运行一次，而控制器在每个游戏步都执行动作。

（第3.2节）。第三，我们用训练好的规划器替换该编码器，并利用动作目标函数联合调整控制器和规划器中的LoRA适配器（第3.3节）。在控制步 t，控制器观测一帧 $o_t \in \mathbb{R}^{H\times W\times 3}$ 并预测

$$a_t = (a_t^{\text{key}}, a_t^{\text{btn}}, a_t^{\text{mouse}}), \tag{1}$$

其中，前两个分量是离散动作，最后一个分量是由鼠标产生的连续相机运动，表示为以度为单位的偏航角和俯仰角变化 $a_t^{\text{mouse}} = (\Delta\psi_t,\ \Delta\theta_t) \in \mathbb{R}^2$。

3.1 任务规划器

视觉锚定分解。任务规划器 $\mathcal{B}$ 将视觉观测和高层指令转换为有序的原子目标序列。给定观测 o 和高层指令 g，它生成

$$\hat{\mathcal{S}} = (\hat{s}_1, \ldots, \hat{s}_{\hat{k}}) \sim \pi_\theta(\cdot \mid o, g), \tag{2}$$

其中，每个 $\hat{s}_j$ 都是一条原子指令。我们使用 GameHorizon Suite（Wang 等人，2026年）的分层游戏过程标注来训练规划器，这些标注将高层目标与时间上对齐的原子指令关联起来，并标注它们之间的先后依赖关系。在训练期间，目标数量 k 通过固定的推理前缀提供，从而得到 $\pi_\theta(\cdot \mid o, g, k)$；推理时则不提供该数量。

我们以 Qwen3.5-9B 为基础，冻结视觉编码器，并使用参考输出序列 $\mathbf{y}$ 训练语言模型和多模态连接器，其中每个目标都以一个目标词元结束（公式（11））：

$$\mathcal{L}_{\text{SFT}} = -\sum_{t=1}^{|\mathbf{y}|} \log \pi_\theta(y_t \mid o, g, k, y_{<t}). \tag{3}$$

此阶段教会规划器以一致的分解格式描述原子目标，并为后续优化提供初始策略。数据统计、超参数以及对朴素监督微调为何会过拟合标注措辞的分析，详见第 A 节。

Dependency-aware refinement. A reference plan records one execution sequence, although multiple orderings may be valid. We refine the SFT policy with GRPO (Shao et al., 2024; DeepSeek-AI, 2025) using separate rewards for step-count agreement and dependency-aware ordering.

For each input, generated steps are aligned one-to-one with reference steps using Hungarian assignment. Pairwise similarity combines content-word F1 and a character-level matching ratio:

$$\mathrm{sim}(a,b) = 0.5\,\mathrm{F1}_{\mathrm{content}}(a,b) + 0.5\ \mathrm{ratio}_{\mathrm{char}}(a,b). \tag{4}$$

Let $\mathrm{pos}[i]$ denote the generated position assigned to reference step i. The count reward is

$$R_k = \max\left(0, 1 - \frac{|\hat{k}-k|}{k}\right). \tag{5}$$

We distinguish two ordering scores. *Dependency satisfaction*, S_{dep}, measures the fraction of annotated must-precede relations respected by the generated plan. *Conventional-order agreement*, S_{conv}, measures agreement with the reference order among pairs without an annotated dependency. The latter expresses a preference for the recorded sequence rather than a requirement for task feasibility.

Let D be the set of must-precede dependencies annotated in GameHorizon Suite (Wang et al., 2026) for a plan and

$$F = \{(i,j) : i < j,\ (i,j) \notin D,\ (j,i) \notin D\}. \tag{6}$$

For a complete alignment,

$$S_{\mathrm{dep}} = \frac{1}{|D|} \sum_{(i,j)\in D} \mathbf{1}[\mathrm{pos}[i] < \mathrm{pos}[j]], \tag{7}$$

$$S_{\mathrm{conv}} = \frac{1}{|F|} \sum_{(i,j)\in F} \mathbf{1}[\mathrm{pos}[i] < \mathrm{pos}[j]]. \tag{8}$$

We combine these scores as

$$R_{\mathrm{ord}} = \mathbf{1}[\hat{k}=k]\,(0.8S_{\mathrm{dep}} + 0.2S_{\mathrm{conv}})\,, \tag{9}$$

$$R = 0.3R_k + 0.7R_{\mathrm{ord}}. \tag{10}$$

The soft count reward distinguishes different degrees of length mismatch. The order reward requires the correct step count and a complete alignment, and prioritizes necessary dependencies over conventional ordering. When one ordering component is undefined, we use the available component alone. GRPO uses mean-centered group rewards and a KL penalty to the SFT reference policy. The 0.8/0.2 split is derived from annotation statistics in Section A.4.

Goal representations. Each atomic goal in the decomposition is terminated by a special goal token τ, so the output sequence is

$$\mathbf{y} = [\mathrm{tok}(\hat{s}_1), \tau, \ldots, \mathrm{tok}(\hat{s}_{\hat{k}}), \tau]. \tag{11}$$

Let p_j be the position of the j-th goal token. The hidden state at this position provides a continuous representation of goal $\hat{s}_j$,

$$z_j = H_\theta(o, g, \mathbf{y})_{p_j}, \tag{12}$$

which summarizes the goal in the context of the observation, the high-level instruction, and the preceding steps. The decoded language supports decomposition supervision and evaluation, while the controller consumes only z_j.

3.2 Frame Controller

The frame controller is a per-frame autoregressive policy that emits native keyboard-and-mouse actions from visual observations. We build on the policy-transformer design of Yue et al. (2026), in which a causal Transformer reads a per-frame token group and produces a single *action-output* token that is decoded into keyboard and mouse actions. We make two changes: the controller is conditioned on the planner's vision-language goal embedding in addition to the visual observation, and continuous camera motion is modeled with a flow-matching head for smooth control.

依赖感知优化。参考计划记录了一种执行序列，但可能存在多种有效排序。我们采用组相对策略优化（Shao 等，2024年；DeepSeek-AI，2025年）优化监督微调策略，分别针对步骤数量一致性和依赖感知排序设置奖励。

对于每个输入，使用匈牙利指派将生成步骤与参考步骤进行一一对齐。成对相似度结合了实词F1分数和字符级匹配率：

$$\mathrm{sim}(a,b) = 0.5\,\mathrm{F1}_{\mathrm{content}}(a,b) + 0.5\ \mathrm{ratio}_{\mathrm{char}}(a,b). \tag{4}$$

Let $\mathrm{pos}[i]$表示分配给参考步骤 i的生成步骤位置。数量奖励为

$$R_k = \max\left(0, 1 - \frac{|\hat{k}-k|}{k}\right). \tag{5}$$

我们区分两种排序得分。依赖满足度 S_{dep} 衡量生成计划遵守的已标注必须先于关系所占的比例。惯常顺序一致性 S_{conv} 衡量无已标注依赖关系的步骤对与参考顺序的一致程度。后者表达的是对所记录顺序的偏好，而非对任务可行性的要求。

设 D 为 GameHorizon Suite（Wang 等人，2026年）中为某个计划标注的必须先于依赖的集合，且

$$F = \{(i,j) : i < j,\ (i,j) \notin D,\ (j,i) \notin D\}. \tag{6}$$

对于完整对齐，

$$S_{\mathrm{dep}} = \frac{1}{|D|} \sum_{(i,j)\in D} \mathbf{1}[\mathrm{pos}[i] < \mathrm{pos}[j]], \tag{7}$$

$$S_{\mathrm{conv}} = \frac{1}{|F|} \sum_{(i,j)\in F} \mathbf{1}[\mathrm{pos}[i] < \mathrm{pos}[j]]. \tag{8}$$

我们将这些得分组合为

$$R_{\mathrm{ord}} = \mathbf{1}[\hat{k}=k]\,(0.8S_{\mathrm{dep}} + 0.2S_{\mathrm{conv}})\,, \tag{9}$$

$$R = 0.3R_k + 0.7R_{\mathrm{ord}}. \tag{10}$$

软数量奖励区分不同程度的长度不匹配。顺序奖励要求正确的步骤数量和完整对齐，并优先考虑必要的依赖关系，而非惯常顺序。当某个顺序分量未定义时，我们仅使用可用的分量。组相对策略优化使用均值中心化组奖励，以及相对于监督微调参考策略的KL惩罚。0.8/0.2的划分源自第A.4节中的标注统计。

目标表征。分解中的每个原子目标均以特殊的目标词元 τ结尾，因此输出序列为

$$\mathbf{y} = [\mathrm{tok}(\hat{s}_1), \tau, \ldots, \mathrm{tok}(\hat{s}_{\hat{k}}), \tau]. \tag{11}$$

设 p_j 为第 j个目标词元的位置。该位置的隐藏状态提供了目标$\hat{s}_j$的连续表示，

$$z_j = H_\theta(o, g, \mathbf{y})_{p_j}, \tag{12}$$

该表示在观测、高层指令和先前步骤的上下文中概括了该目标。解码语言支持分解监督和评估，而控制器仅使用z_j。

3.2 帧控制器

帧控制器是一种逐帧自回归策略，可根据视觉观测输出原生键鼠动作。我们基于Yue等人（2026年）的策略Transformer设计构建该控制器，其中，因果Transformer读取逐帧词元组，并生成单个动作输出词元，再将其解码为键盘和鼠标动作。我们进行了两项改动：除视觉观测外，控制器还以规划器的视觉-语言目标嵌入为条件；同时，使用流匹配头对连续相机运动进行建模，以实现平滑控制。

Frame-level conditioning. Training samples are windows of T consecutive frames, starting at t_0, that lie within one annotated atomic segment. At frame t, the controller consumes the goal embedding z (Section 3.3), a visual token o_t obtained by feeding the frame into a frozen SigLIP encoder (Zhai et al., 2023) followed by a Perceiver resampler, a learned action-query token a_{in}, and the tokens of action a_t, which only later frames can attend to. The output of the causal Transformer f_ϕ at the position of a_{in} is the action-output token

$$a_t^{\text{out}} = f_\phi(z,\, o_{t_0:t},\, a_{t_0:t-1};\, a_{\text{in}})\,, \qquad t_0 \le t < t_0 + T, \tag{13}$$

from which the action heads decode $p_\phi(a_t \mid a_t^{\text{out}})$. The previous actions $a_{t_0:t-1}$ are ground truth during training and the controller's own outputs at inference.

Hybrid action head. The action-output token a_t^{out} is decoded by two heads. Discrete keyboard keys and mouse buttons are produced by a small autoregressive action decoder that samples the action tokens one at a time from a_t^{out} (Yue et al., 2026). Continuous camera motion is instead generated by a lightweight flow-matching head (Lipman et al., 2023; Liu et al., 2023), which avoids discretizing small corrections and can represent several plausible directions. The action objective is

$$\mathcal{L}_{\text{act}} = \mathcal{L}_{\text{key}} + \mathcal{L}_{\text{btn}} + \lambda_{\text{mouse}}\mathcal{L}_{\text{mouse}}, \tag{14}$$

where $\mathcal{L}_{\text{key}}$ and $\mathcal{L}_{\text{btn}}$ are the cross-entropy losses of the autoregressive decoder and $\mathcal{L}_{\text{mouse}}$ is the flow-matching loss.

Action conditioning and causal masking. Following Yue et al. (2026), the action-output token at frame t cannot attend to the same-frame action tokens, which would leak the target, and no token attends to the action-output token; all other tokens are causal (Figure 5).

3.3 Planner–Controller Alignment

During controller pretraining (Stage 2), the goal embedding z_0 is produced by a frozen, generic vision-language model $\mathcal{E}_0$ (Qwen3.5-9B) from the frame at the start of the window and its atomic instruction s. In the third stage, we replace $\mathcal{E}_0$ with the GRPO-trained planner $\mathcal{B}$, augmented with LoRA adapters (Hu et al., 2022) Δ in every Transformer block of its language model. Each training window provides one annotated atomic goal s, so the planner reads the first frame of the window and s, followed by the goal token, and the hidden state at that token becomes the condition,

$$z_{\mathcal{B}} = H_{\mathcal{B}+\Delta}([\,o_{t_0}, \text{tok}(s), \tau\,])_\tau\,, \tag{15}$$

which replaces z_0 as the goal embedding z in Equation (13). The objective in Equation (14) updates all controller parameters and Δ while the base planner weights remain frozen. The resulting adapted planner is the single planner used at inference, for both decomposition and goal embeddings. Additional training details of this stage are given in Section B.3.

3.4 Multi-Rate Inference

Inference runs at two rates. *Task rate:* for a high-level goal g, the adapted planner decodes the plan $(\hat{s}_1, \ldots, \hat{s}_{\hat{k}})$ once, without a supplied step count, emitting a goal token after every atomic goal; the hidden states at these tokens are cached as $(z_1, \ldots, z_{\hat{k}})$ (Equation (12)). *Frame rate:* the controller receives only the cached embedding of the active goal, z_j, combines it with the frames and actions of the current window, and emits an action at every game step. The decoded text is available for inspection and evaluation but is not passed to the controller.

4 Experiments

We ask three questions: 1) Does TaskFrame act well in closed loop (Section 4.2)? 2) Does its planner produce good plans (Section 4.3)? 3) How should planning and control be connected (Section 4.4)?

帧级条件控制。训练样本是包含 T 个连续帧的窗口，起始于 t_0，且均位于同一个已标注原子片段内。在帧 t 处，控制器接收目标嵌入 z （第3.3节）、视觉词元 o_t （通过将该帧输入冻结的SigLIP编码器（Zhai等人，2023年），再经过Perceiver重采样器获得）、可学习的动作查询词元 a_{in}，以及动作 a_t的词元，这些动作词元仅可被后续帧关注。因果Transformer f_ϕ 在 a_{in} 位置处的输出即为动作输出词元

$$a_t^{\text{out}} = f_\phi(z,\, o_{t_0:t},\, a_{t_0:t-1};\, a_{\text{in}})\,, \qquad t_0 \le t < t_0 + T, \tag{13}$$

动作头从中解码出 p_ϕ （$a_t \mid a_t^{\text{out}}$）。先前的动作$a_{t_0:t-1}$ 在训练时为真实标注，在推理时则为控制器自身的输出。

混合动作头。动作输出词元 a_t^{out}由两个注意力头解码。离散的键盘按键和鼠标按钮由一个小型自回归动作解码器生成，该解码器从 a_t^{out}中逐个采样动作词元（Yue等人，2026年）。连续相机运动则由一个轻量级流匹配头生成（Lipman等人，2023年；Liu等人，2023年），这避免了对微小修正进行离散化，并能表示多个合理的方向。动作目标函数为

$$\mathcal{L}_{\text{act}} = \mathcal{L}_{\text{key}} + \mathcal{L}_{\text{btn}} + \lambda_{\text{mouse}}\mathcal{L}_{\text{mouse}}, \tag{14}$$

其中， $\mathcal{L}_{\text{key}}$ 和 $\mathcal{L}_{\text{btn}}$ 是自回归解码器的交叉熵损失， $\mathcal{L}_{\text{mouse}}$ 是流匹配损失。

动作条件控制与因果掩码。遵循Yue等人（2026年）的方法，第 t帧的动作输出词元不能关注同帧动作词元，否则会泄露目标；同时，任何词元都不能关注动作输出词元；其余所有词元均采用因果注意力（图5）。

3.3 规划器与控制器对齐

在控制器预训练（第二阶段）期间，目标嵌入 z_0 由冻结的通用视觉语言模型 $\mathcal{E}_0$ （Qwen3.5-9B）根据窗口起始帧及其原子指令 s生成。在第三阶段，我们将 $\mathcal{E}_0$ 替换为经GRPO训练的规划器 $\mathcal{B}$，并在其语言模型的每个Transformer块中加入LoRA适配器（Hu等人，2022年） Δ。每个训练窗口提供一个带标注的原子目标s，因此规划器读取窗口的第一帧和 s，随后读取目标词元，并将该词元处的隐藏状态作为条件，

$$z_{\mathcal{B}} = H_{\mathcal{B}+\Delta}([\,o_{t_0}, \text{tok}(s), \tau\,])_\tau\,, \tag{15}$$

该条件替代 z_0 ，作为目标嵌入 z 用于公式（13）。公式（14）中的目标函数用于更新所有控制器参数和 Δ ，而基础规划器权重保持冻结。由此得到的适配后的规划器是推理时唯一使用的规划器，同时用于任务分解和生成目标嵌入。本阶段的更多训练细节见第B.3节。

3.4 多速率推理

推理以两种速率运行。任务级速率：对于高层目标 g，适配后的规划器仅解码一次计划 $(\hat{s}_1, \ldots, \hat{s}_{\hat{k}})$，无需提供步骤数量，并在每个原子目标之后输出一个目标词元；这些词元处的隐藏状态被缓存为 $(z_1, \ldots, z_{\hat{k}})$ （公式(12)）。帧率：控制器仅接收当前目标的缓存嵌入 z_j，将其与当前窗口内的帧和动作相结合，并在每个游戏步输出一个动作。解码文本可用于检查和评估，但不会传递给控制器。

4 实验

我们提出三个问题：1）TaskFrame 在闭环中的表现是否良好（第 4.2 节）？ 2）其规划器能否生成良好的计划（第 4.3 节）？ 3）应如何连接规划与控制（第 4.4 节）？

Table 1 Closed-loop success rate (%) on atomic Minecraft tasks over 3 evaluation seeds, with 12 episodes per kill task and 6 per mining task in each seed. Averages are reported with a binomial estimate of their standard deviation across seeds; per-task mining results are in Section C.3.

Model	Kill Entities					Mine Blocks
	Cow	Sheep	Pig	Zombie	Avg.	Avg. of 4
VPT-BC (Baker et al., 2022)	17	17	17	8	15 ±5.4	38 ±9.2
VPT-RL (Baker et al., 2022)	8	17	8	0	8 ±4.5	29 ±8.8
STEVE-1 (Lifshitz et al., 2023)	17	8	17	8	13 ±5.2	62 ±9.2
JARVIS-VLA-7B (Li et al., 2025a)	17	33	**50**	25	31 ±6.6	**100** ±5.3
TaskFrame (Ours)	**67**	**58**	42	**35**	**51** ±7.0	96 ±6.2

4.1 Setup

Training. The controller is trained on about 100 hours of VPT contractor gameplay (Baker et al., 2022) with native keyboard-and-mouse actions, using the atomic instructions of GameHorizon Suite (Wang et al., 2026) as conditions. Training windows contain $T = 21$ frames (about one second at 20 Hz) and never cross an atomic-segment boundary. The backbone is a 30-layer causal Transformer of width 1024. The planner is trained on the decomposition annotations of GameHorizon Suite.

Closed-loop evaluation. We follow the programmatic protocol of JARVIS-VLA (Li et al., 2025a) and MCU (Zheng et al., 2025) in MineStudio (Cai et al., 2024a) and evaluate two types of tasks: *mine* tasks place the equipped agent inside a chamber of the target block, isolating block breaking from search, and *kill* tasks spawn target animals a few blocks ahead of an agent holding a sword. Episodes last at most 400 steps. We compare with VPT-BC, VPT-RL (Baker et al., 2022), STEVE-1 (Lifshitz et al., 2023), and JARVIS-VLA-7B (Li et al., 2025a). Task configurations and baseline settings are given in Section C.

Planner evaluation. *Soft Semantic F1* We utilize (SSF1) to measure step content on 2.1K held-out plans. The planner sees one frame and the high-level instruction without the step count, and the predicted and reference steps are matched by sentence-embedding similarity, so paraphrases are accepted while missing or redundant steps are penalized. *Task2* of GameHorizon Suite (Wang et al., 2026) measures step ordering on 250 questions: given the step count, the planner generates a plan, which is matched to the most similar of four candidate orderings; ties and unparseable plans count as errors. More details are provided in Section A.2.

4.2 Closed-Loop Control

Main results. Table 1 reports the results. TaskFrame reaches an average kill success of 51%, compared with 31% for JARVIS-VLA-7B, and performs best on cow, sheep, and zombie. This gap is about twice the combined standard deviation. On mining, both TaskFrame and JARVIS-VLA-7B succeed in almost every episode, so this category is close to saturated in our setting. VPT-BC, VPT-RL, and STEVE-1 reach 8–15% on kill tasks and 29–62% on mining. Overall, TaskFrame achieves the best combat performance among the evaluated models while remaining on par with JARVIS-VLA-7B on mining.

Qualitative results. Figure 3 shows closed-loop rollouts. In combat, the controller turns toward the target animals, approaches them through the grass, and keeps attacking until the target is killed, even when several animals of different types are in view; in mining, it repeatedly breaks the log in front of it.

4.3 Task Planning

As shown in Table 2, our 9B planner reaches an SSF1 of 0.3459 and a Task2 accuracy of 40.0%, exceeding GPT-4o and Qwen3.5-397B-A17B on both metrics while remaining below the two Gemini models, which are evaluated zero-shot like the other general-purpose baselines. Over the backbone, SFT raises SSF1 from 0.2874

表 1 在 3 个评估随机种子下，原子级《我的世界》任务的闭环成功率（%）。每个随机种子下，每项击杀任务进行 12 个回合，每项采矿任务进行 6 个回合。报告平均值及其跨随机种子标准差的二项分布估计；各项采矿任务的结果见第 C.3 节。

模型	击杀实体					挖掘方块
	Cow	羊	Pig	僵尸	Avg.	4项平均值
VPT-BC（Baker 等人，2022年）	17	17	17	8	15 ±5.4	38 ±9.2
VPT-RL（Baker 等人，2022年）	8	17	8	0	8 ±4.5	29 ±8.8
STEVE-1（Lifshitz 等人，2023年）	17	8	17	8	13 ±5.2	62 ±9.2
JARVIS-VLA-7B（Li 等人，2025a）	17	33	**50**	25	31 ±6.6	**100** ±5.3
TaskFrame（我们的方法）	**67**	**58**	42	**35**	**51** ±7.0	96 ±6.2

4.1 实验设置

训练。控制器使用约 100 小时包含原生键鼠动作的 VPT 外包人员游戏操作数据（Baker 等人，2022年）进行训练，并以 GameHorizon Suite（Wang 等人，2026年）的原子指令作为条件。训练窗口包含 $T = 21$ 帧（在 20 Hz 下约为一秒），且从不跨越原子片段边界。骨干网络是一个宽度为 1024 的 30 层因果Transformer。规划器使用 GameHorizon Suite 的分解标注进行训练。

闭环评估。我们在 MineStudio 中遵循 JARVIS-VLA（Li 等人，2025a）和 MCU（Zheng 等人，2025年）的程序化协议（MineStudio：Cai 等人，2024a），评估两类任务：挖掘任务将配备好装备的智能体置于由目标方块构成的密室内，将方块破坏与搜索分离；击杀任务则在持剑智能体前方几个方块处生成目标动物。每个回合最多持续 400 步。我们与 VPT-BC、VPT-RL（Baker 等人，2022年）、STEVE-1（Lifshitz 等人，2023年）以及 JARVIS-VLA-7B（Li 等人，2025a）进行比较。任务配置和基线设置见附录 C 节。

规划器评估。软语义 *F1* 我们使用软语义 F1（SSF1）来衡量 2.1K 个留出计划的步骤内容。规划器接收一帧画面和不含步骤数量的高层指令，并通过句子嵌入相似度匹配预测步骤与参考步骤，因此允许语义相同的不同表述，但会对缺失或冗余步骤进行惩罚。*Task2* 是 GameHorizon Suite（Wang 等人，2026年）中的一项任务，通过 250 道问题衡量步骤排序：给定步骤数量后，规划器生成一个计划，并将其与四种候选排序中最相似的一种进行匹配；出现并列匹配或计划无法解析时均计为错误。更多细节见第 A.2 节。

4.2 闭环控制

主要结果。表 1 报告了实验结果。TaskFrame 的平均击杀成功率达到 51%，而 JARVIS-VLA-7B 为 31%，且在牛、羊和僵尸上的表现最佳。这一差距约为合并标准差的两倍。在采矿任务中，TaskFrame 和 JARVIS-VLA-7B 几乎每个回合都能成功，因此在我们的设置下，该类别的性能已接近饱和。VPT-BC、VPT-RL 和 STEVE-1 在击杀任务中的成功率为 8–15%，在采矿任务中为 29–62%。总体而言，TaskFrame 在受评估模型中取得了最佳战斗性能，同时在采矿任务上与 JARVIS-VLA-7B 保持相当的水平。

定性结果。图 3展示了闭环轨迹展开。在战斗中，控制器转向目标动物，穿过草丛接近它们，并持续攻击直至目标被击杀，即使视野中同时出现多只不同种类的动物也不例外；在采矿中，控制器反复破坏面前的原木。

4.3 任务规划

如表2所示，我们的90亿参数规划器的SSF1达到0.3459，Task2准确率达到40.0%，两项指标均超过GPT-4o和Qwen3.5-397B-A17B，但仍低于两个Gemini模型；这两个模型与其他通用基线一样，均采用零样本评估。相较于骨干网络，监督微调将SSF1从0.2874提高

Figure 3 Closed-loop rollouts of TaskFrame. Each row shows four frames of one successful episode on an atomic task. Comparisons with baselines are shown in Figure 6.

to 0.3312 and Task2 accuracy from 23.2% to 38.8%, and GRPO further improves both metrics and lowers the Unparsed rate from 14.0% to 11.6%. The gain in SSF1, which uses sentence embeddings rather than the reward's step matcher and supplies no step count, indicates that GRPO does not merely fit its own reward.

4.4 Connecting Planning and Control

Planner and controller roles. To isolate the value of the hierarchy and of joint alignment, we predict actions on held-out gameplay with high-level goals (Table 3; setup in Section C.6). *Planner only* generates actions directly with the planner; *controller only* uses an empty goal; *cascade* passes the planner's decoded instructions as text to the Stage-2 controller; TaskFrame passes the aligned planner's goal embeddings; and *oracle goals* uses annotated atomic instructions as an upper bound.

The planner alone beats the unconditioned controller, but it emits actions for a whole window from a single observation and cannot react within the window; the controller alone cannot tell which behavior is intended. The cascade recovers most of the gap, and the aligned goal embeddings of TaskFrame give the best results on all metrics; since both use the same decomposition, the gap between them isolates the effect of joint alignment. Oracle goals add little, so most remaining error lies in the controller rather than in the decomposition.

Goal conditioning in closed loop. We keep the trained model fixed and change only the goal it receives, evaluating each setting on the kill tasks with 48 episodes per task in each of three seeds (Table 4; per-task results in Section C.5). With an empty goal, kill success drops from 51% to 25%; some success remains because every spawned animal is a valid target. A mismatched goal lowers it further to 13%, so the controller actively follows the goal rather than ignoring it, and withholding the observation from the planner costs 20 points, showing that the visual grounding of the goal helps.

Camera head. We compare the flow-matching camera head with the discrete head of Yue et al. (2026), which classifies camera motion into bins and takes their expected value. The two controllers differ only in the camera head and are trained with the same data and schedule (Table 5; setup in Section B.4). Flow matching lowers the mean absolute error from 2.1° to 1.7°

Table 5 Camera MAE (°) on held-out gameplay.

Camera head	Yaw ↓	Pitch ↓
Discrete bins	2.1	1.2
Flow matching	**1.7**	**1.0**

图 3 TaskFrame 的闭环轨迹展开。每行展示一个原子任务的成功回合中的四帧。与基线的比较见图 6。

至0.3312，并将Task2准确率从23.2%提高至38.8%；组相对策略优化进一步提升了这两项指标，并将未解析率从14.0%降至11.6%。SSF1使用句子嵌入，而非奖励所采用的步骤匹配器，且不提供步骤数量；该指标的提升表明，组相对策略优化并非仅仅拟合自身的奖励。

4.4 连接规划与控制

规划器与控制器的作用。为单独评估层次结构与联合对齐的价值，我们在具有高层目标的留出游戏过程上预测动作（表 3；实验设置见第 C.6 节）。仅使用规划器时，直接由规划器生成动作；仅使用控制器时，使用空目标；级联将规划器解码得到的指令以文本形式传递给第二阶段控制器；TaskFrame 则传递对齐后的规划器的目标嵌入；而预言目标使用标注的原子指令作为上界。

单独使用规划器优于无条件控制器，但它仅根据单次观测就输出整个窗口的动作，无法在窗口内作出反应；单独使用控制器则无法辨别预期行为。级联弥补了大部分差距，而 TaskFrame 对齐后的目标嵌入在所有指标上均取得最佳结果；由于二者采用相同的分解方式，它们之间的差距单独体现了联合对齐的效果。预言目标带来的提升很小，因此剩余误差大多源于控制器，而非分解方式。

闭环中的目标条件化。我们保持训练好的模型不变，仅改变其接收的目标，在击杀任务上评估各项设置：使用三个随机种子，每个随机种子下每项任务进行48个回合（表4；各任务的结果见第C.5节）。使用空目标时，击杀成功率从51%降至25%；由于每只生成的动物都是有效目标，因此仍能取得一定的成功。不匹配目标使击杀成功率进一步降至13%，说明控制器会主动遵循目标，而非忽略目标；不向规划器提供观测则会使成功率下降20个百分点，表明目标的视觉锚定有所帮助。

相机预测头。我们将流匹配相机预测头与Yue等人（2026年）的离散预测头进行比较，后者将相机运动划分到不同分箱中，并取其期望值。这两个控制器仅在相机预测头上有所不同，且使用相同的数据和学习率调度进行训练（表5；设置见第B.4节）。流匹配将平均绝对误差从2.1° 降至1.7°

表5 在留出的游戏过程数据上的相机平均绝对误差（°）。

相机预测头	偏航角 ↓	俯仰角 ↓
离散分箱	2.1	1.2
流匹配	**1.7**	**1.0**

Table 2 Planner performance. SSF1 is evaluated without a supplied step count, whereas Task2 provides the target count. Unparsed is the percentage of questions without a valid decision.

Model	SSF1 ↑	Task2	
		Acc. (%) ↑	Unparsed (%) ↓
Gemini 3.5 Flash (Google DeepMind, 2026a)	**0.3815**	44.4	13.6
Gemini 3.6 Flash (Google DeepMind, 2026b)	0.3808	**46.8**	15.2
GPT-4o (OpenAI, 2024)	0.2720	36.0	14.0
Qwen3-Omni-30B-A3B (Qwen Team, 2025)	0.2643	34.0	11.6
Qwen3.5-397B-A17B (Qwen Team, 2026a)	0.2990	35.6	12.4
Qwen3.6-35B-A3B (Qwen Team, 2026b)	0.2951	34.4	**8.4**
BAGEL-7B-MoT (Deng et al., 2025)	–	28.0	14.4
SenseNova-U1-8B-MoT (Diao et al., 2026)	–	27.6	14.4
Qwen3.5-9B (backbone)	0.2874	23.2	23.2
+ SFT	0.3312	38.8	14.0
+ GRPO (**TaskFrame planner**)	0.3459	40.0	11.6

Table 3 Offline action prediction on held-out gameplay with high-level goals.

Method	Key F1 ↑	Btn. F1 ↑	Cam. RMSE ↓
Planner only	0.21	0.38	6.6
Controller only	0.16	0.32	6.8
Cascade (text)	0.24	0.45	5.9
TaskFrame	**0.27**	**0.49**	**5.6**
Oracle goals	0.29	0.52	5.4

Table 4 Closed-loop kill success (%) when only the supplied goal is changed.

Goal	Kill avg. ↑
TaskFrame (full)	**51** ±3.5
Empty goal	25 ±3.0
Mismatched goal	13 ±2.3
Planner w/o observation	31 ±3.1

for yaw and from 1.2° to 1.0° for pitch. Binning rounds small corrections to the nearest bin and averages over several plausible directions, which hurts most when aiming at a moving target.

5 Conclusion

We presented TaskFrame, a hierarchical game agent that separates task-level planning from frame-level keyboard-and-mouse control. A vision-language planner, trained with supervised decomposition and dependency-aware GRPO, turns a high-level instruction into an ordered sequence of atomic goals and outperforms general-purpose models such as GPT-4o and Qwen3.5-397B-A17B on instruction decomposition. The hidden states of the goal tokens in its plan serve as goal embeddings. A causal Transformer controller with a hybrid action head converts the active embedding and recent frames into native actions at every game step. In closed-loop Minecraft evaluation, the agent achieves the highest combat success rate among all evaluated models. Ablations further show that the controller relies on the goal condition and that the aligned goal embeddings are more effective than passing decoded text between the two modules.

Limitations. The agent does not decide on its own when an atomic goal is done, so it relies on an external execution protocol to switch to the next goal, and it does not replan after a failure. In addition, the controller sees only about one second of history, which is short for long-horizon tasks. Finally, all experiments are conducted in Minecraft, and extending the method to other games remains a challenging direction for future work.

表 2 规划器性能。SSF1 在未提供步骤数量的情况下进行评估，而 Task2 则提供目标数量。未解析比例指未获得有效决策的问题所占的百分比。

模型	SSF1 ↑	Task2	
		准确率（%）↑	未解析比例（%）↓
Gemini 3.5Flash（谷歌 DeepMind,2026a）	**0.3815**	44.4	13.6
Gemini 3.6 Flash（谷歌 DeepMind，2026b）	0.3808	**46.8**	15.2
GPT-4o（OpenAI，2024年）	0.2720	36.0	14.0
Qwen3-Omni-30B-A3B（通义千问团队，2025年）	0.2643	34.0	11.6
Qwen3.5-397B-A17B（通义千问团队,2026a）	0.2990	35.6	12.4
Qwen3.6-35B-A3B（通义千问团队,2026b）	0.2951	34.4	**8.4**
BAGEL-7B-MoT（Deng等人，2025年）	–	28.0	14.4
SenseNova-U1-8B-MoT（Diao 等，2026年）	–	27.6	14.4
Qwen3.5-9B（骨干网络）	0.2874	23.2	23.2
+ 监督微调	0.3312	38.8	14.0
+ 组相对策略优化（**TaskFrame 规划器**）	0.3459	40.0	11.6

表 3 在高层目标引导下对留出游戏过程进行离线动作预测。

方法	按键 F1 ↑	按钮 F1 ↑	视角均方根误差 ↓
仅规划器	0.21	0.38	6.6
仅控制器	0.16	0.32	6.8
级联（文本）	0.24	0.45	5.9
TaskFrame	**0.27**	**0.49**	**5.6**
预言目标	0.29	0.52	5.4

表 4 仅改变所提供目标时的闭环击杀成功率（%）。

Goal	平均击杀成功率 ↑
TaskFrame（完整版）	**51** ±3.5
空目标	25 ±3.0
不匹配目标	13 ±2.3
无观测规划器	31 ±3.1

对于偏航角，而俯仰角则从 1.2° 变为 1.0° 。分箱会将微小修正舍入到最近的箱，并对多个可能的方向取平均，这在瞄准移动目标时影响最大。

5 结论

我们提出了TaskFrame，这是一种将任务级规划与帧级键鼠控制分离的分层游戏智能体。视觉语言规划器通过监督式分解和依赖感知GRPO进行训练，将高层指令转化为有序的原子目标序列，并在指令分解方面优于GPT-4o和Qwen3.5-397B-A17B等通用模型。其规划中目标词元的隐藏状态用作目标嵌入。带有混合动作头的因果Transformer控制器在每个游戏步将当前激活的嵌入和最近的帧转换为原生动作。在Minecraft闭环评估中，该智能体在所有参评模型中取得了最高的战斗成功率。消融实验进一步表明，控制器依赖于目标条件，而且与在两个模块之间传递解码文本相比，对齐后的目标嵌入更为有效。

局限性。智能体无法自主判断原子目标何时完成，因此依赖外部执行协议来切换至下一个目标，且在失败后不会重新规划。此外，控制器只能看到约一秒的历史信息，对于长时程任务而言，这一时间范围较短。最后，所有实验均在我的世界中开展，将该方法扩展到其他游戏仍是未来工作中一个具有挑战性的方向。

References

Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Chuyuan Fu, Keerthana Gopalakrishnan, Karol Hausman, et al. Do as i can, not as i say: Grounding language in robotic affordances. *arXiv preprint arXiv:2204.01691*, 2022.

Adrià Puigdomènech Badia, Bilal Piot, Steven Kapturowski, Pablo Sprechmann, Alex Vitvitskyi, Zhaohan Daniel Guo, and Charles Blundell. Agent57: Outperforming the Atari human benchmark. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 507–517. PMLR, 2020. https://proceedings.mlr.press/v119/badia20a.html.

Bowen Baker, Ilge Akkaya, Peter Zhokhov, Joost Huizinga, Jie Tang, Adrien Ecoffet, Brandon Houghton, Raul Sampedro, and Jeff Clune. Video pretraining (vpt): Learning to act by watching unlabeled online videos. In *Advances in neural information processing systems*, volume 35, pages 24639–24654, 2022.

Suneel Belkhale, Tianli Ding, Ted Xiao, Pierre Sermanet, Quan Vuong, Jonathan Tompson, Yevgen Chebotar, Debidatta Dwibedi, and Dorsa Sadigh. RT-H: Action hierarchies using language. *arXiv preprint arXiv:2403.01823*, 2024.

Kevin Black, Noah Brown, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Lachy Groom, Karol Hausman, Brian Ichter, et al. π_0: A vision-language-action flow model for general robot control. *arXiv preprint arXiv:2410.24164*, 2024.

Kevin Black, Manuel Y. Galliker, and Sergey Levine. Real-time execution of action chunking flow policies. *arXiv preprint arXiv:2506.07339*, 2025.

Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Xi Chen, Krzysztof Choromanski, Tianli Ding, Danny Driess, Avinava Dubey, Chelsea Finn, et al. Rt-2: Vision-language-action models transfer web knowledge to robotic control. *arXiv preprint arXiv:2307.15818*, 2023.

Noam Brown and Tuomas Sandholm. Superhuman AI for multiplayer poker. *Science*, 365(6456):885–890, 2019. doi: 10.1126/science.aay2400.

Qingwen Bu, Hongyang Li, Li Chen, Jisong Cai, Jia Zeng, Heming Cui, Maoqing Yao, and Yu Qiao. Towards synergistic, generalized and efficient dual-system for robotic manipulation. *arXiv preprint arXiv:2410.08001*, 2024.

Shaofei Cai, Zhancun Mu, Kaichen He, Bowei Zhang, Xinyue Zheng, Anji Liu, and Yitao Liang. MineStudio: A streamlined package for Minecraft AI agent development. *arXiv preprint arXiv:2412.18293*, 2024a.

Shaofei Cai, Bowei Zhang, Zihao Wang, Xiaojian Ma, Anji Liu, and Yitao Liang. Groot: Learning to follow instructions by watching gameplay videos. In *International Conference on Learning Representations*, 2024b.

Shaofei Cai, Zihao Wang, Kewei Lian, Zhancun Mu, Xiaojian Ma, Anji Liu, and Yitao Liang. Rocket-1: Mastering open-world interaction with visual-temporal context prompting. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 12122–12131, 2025.

Cheng Chi, Zhenjia Xu, Siyuan Feng, Eric Cousineau, Yilun Du, Benjamin Burchfiel, Russ Tedrake, and Shuran Song. Diffusion policy: Visuomotor policy learning via action diffusion. *The International Journal of Robotics Research*, 44(10-11):1684–1704, 2025.

DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

Chaorui Deng, Deyao Zhu, Kunchang Li, Chenhui Gou, Feng Li, Zeyu Wang, Shu Zhong, Weihao Yu, Xiaonan Nie, Ziang Song, et al. Emerging properties in unified multimodal pretraining. *arXiv preprint arXiv:2505.14683*, 2025.

Haiwen Diao, Penghao Wu, Hanming Deng, Jiahao Wang, Shihao Bai, Silei Wu, Weichen Fan, Wenjie Ye, Wenwen Tong, Xiangyu Fan, et al. Sensenova-u1: Unifying multimodal understanding and generation with NEO-unify architecture. *arXiv preprint arXiv:2605.12500*, 2026.

Linxi Fan, Guanzhi Wang, Yunfan Jiang, Ajay Mandlekar, Yuncong Yang, Haoyi Zhu, Andrew Tang, De-An Huang, Yuke Zhu, and Anima Anandkumar. Minedojo: Building open-ended embodied agents with internet-scale knowledge. In *Advances in neural information processing systems*, volume 35, pages 18343–18362, 2022.

Figure AI. Helix: A vision-language-action model for generalist humanoid control. Technical report, 2025. https://www.figure.ai/news/helix.

参考文献

Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Chuyuan Fu, Keerthana Gopalakrishnan, Karol Hausman 等。按我所能，而非按我所言：将语言扎根于机器人可供性。arXiv 预印本 arXiv:2204.01691，2022年。Adrià Puigdomènech Badia, Bilal Piot, Steven Kapturowski, Pablo Sprechmann, Alex Vitvitskyi, Zhaohan Daniel Guo 和 Charles Blundell。Agent57：超越 Atari人类基准。载于第37届国际机器学习大会论文集，机器学习研究论文集第119卷，第507–517页。PMLR，2020年。https://proceedings.mlr.press/v119/badia20a.html。Bowen Baker, Ilge Akkaya, Peter Zhokhov, Joost Huizinga, Jie Tang, Adrien Ecoffet, Brandon Houghton, Raul Sampedro 和 Jeff Clune。视频预训练（VPT）：通过观看无标签在线视频学习行动。载于神经信息处理系统进展，第35卷，第24639–24654页，2022年。Suneel Belkhale, Tianli Ding, Ted Xiao, Pierre Sermanet, Quan Vuong, Jonathan Tompson, Yevgen Chebotar, Debidatta Dwibedi 和 Dorsa Sadigh。RT-H：利用语言构建动作层级。arXiv 预印本 arXiv:2403 01823，2024年。. Kevin Black, Noah Brown, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Lachy Groom, Karol Hausman, Brian Ichter 等。 π_0：用于通用机器人控制的视觉-语言-动作流模型。arXiv 预印本 arXiv:2410.24164，2024年。Kevin Black, Manuel Y. Galliker 和 Sergey Levine。动作分块流策略的实时执行。arXiv 预印本 arXiv:2506.07339，2025年。Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Xi Chen, Krzysztof Choromanski, Tianli Ding, Danny Driess, Avinava Dubey, Chelsea Finn 等。RT-2：视觉-语言-动作模型将网络知识迁移至机器人控制。arXiv 预印本 arXiv:2307.15818，2023年。Noam Brown 和 Tuomas Sandholm。用于多人扑克的超人类人工智能。科学，365(6456):885–890，2019年。doi: 10.1126/science.aay2400。Qingwen Bu, Hongyang Li, Li Chen, Jisong Cai, Jia Zeng, Heming Cui, Maoqing Yao 和 Yu Qiao。迈向用于机器人操作的协同、通用且高效的双系统。arXiv 预印本 arXiv:2410.08001，2024年。Shaofei Cai, Zhancun Mu, Kaichen He, Bowei Zhang, Xinyue Zheng, Anji Liu 和 Yitao Liang。MineStudio：用于我的世界人工智能智能体开发的精简工具包。arXiv 预印本 arXiv:2412.18293，2024a。Shaofei Cai, Bowei Zhang, Zihao Wang, Xiaojian Ma, Anji Liu 和 Yitao Liang。GROOT：通过观看游戏实况视频学习遵循指令。载于国际学习表征会议，2024b。Shaofei Cai, Zihao Wang, Kewei Lian, Zhancun Mu, Xiaojian Ma, Anji Liu 和 Yitao Liang。ROCKET-1：通过视觉时序上下文提示掌握开放世界交互。载于 IEEE/CVF计算机视觉与模式识别会议（CVPR）论文集，第12122–12131页，2025年。Cheng Chi, Zhenjia Xu, Siyuan Feng, Eric Cousineau, Yilun Du, Benjamin Burchfiel, Russ Tedrake 和 Shuran Song。扩散策略：通过动作扩散进行视觉运动策略学习。国际机器人研究杂志，44(10-11):1684–1704，2025年。DeepSeek-AI。DeepSeek-R1：通过强化学习激发大语言模型的推理能力。arXiv 预印本 arXiv:2501.12948，2025年。Chaorui Deng, Deyao Zhu, Kunchang Li, Chenhui Gou, Feng Li, Zeyu Wang, Shu Zhong, Weihao Yu, Xiaonan Nie, Ziang Song 等。统一多模态预训练中的涌现特性。arXiv 预印本 arXiv:2505.14683，2025年。Haiwen Diao, Penghao Wu, Hanming Deng, Jiahao Wang, Shihao Bai, Silei Wu, Weichen Fan, Wenjie Ye, Wenwen Tong, Xiangyu Fan 等。Sensenova-u1：通过 NEO-unify架构统一多模态理解与生成。arXiv 预印本 arXiv:2605.12500，2026年。Linxi Fan, Guanzhi Wang, Yunfan Jiang, Ajay Mandlekar, Yuncong Yang, Haoyi Zhu, Andrew Tang, De-An Huang, Yuke Zhu 和 Anima Anandkumar。MineDojo：利用互联网规模知识构建开放式具身智能体。载于神经信息处理系统进展，第35卷，第18343–18362页，2022年。Figure AI。Helix：用于通用人形机器人控制的视觉-语言-动作模型。技术报告，2025年。https://www.figure.ai/news/Helix.

Google DeepMind. Gemini 3.5 Flash model card. https://deepmind.google/models/model-cards/gemini-3-5-flash/, May 2026a.

Google DeepMind. Gemini 3.6 Flash model card. https://deepmind.google/models/model-cards/gemini-3-6-flash/, July 2026b.

William H. Guss, Brandon Houghton, Nicholay Topin, Phillip Wang, Cayden Codel, Manuela Veloso, and Ruslan Salakhutdinov. MineRL: A large-scale dataset of Minecraft demonstrations. In *International Joint Conference on Artificial Intelligence*, 2019.

Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse control tasks through world models. *Nature*, 640(8059):647–653, April 2025. doi: 10.1038/s41586-025-08744-2. https://doi.org/10.1038/s41586-025-08744-2.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.

Wenlong Huang, Fei Xia, Ted Xiao, Harris Chan, Jacky Liang, Pete Florence, Andy Zeng, Jonathan Tompson, Igor Mordatch, Yevgen Chebotar, et al. Inner monologue: Embodied reasoning through planning with language models. In *Conference on Robot Learning*, 2022.

Daniel Kahneman. *Thinking, Fast and Slow*. Farrar, Straus and Giroux, 2011.

Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael Rafailov, Ethan Foster, Grace Lam, Pannag Sanketi, et al. Openvla: An open-source vision-language-action model. *arXiv preprint arXiv:2406.09246*, 2024.

Moo Jin Kim, Chelsea Finn, and Percy Liang. Fine-tuning vision-language-action models: Optimizing speed and success. *arXiv preprint arXiv:2502.19645*, 2025.

Muyao Li, Zihao Wang, Kaichen He, Xiaojian Ma, and Yitao Liang. Jarvis-vla: Post-training large-scale vision language models to play visual games with keyboards and mouse. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 17878–17899, 2025a.

Zaijing Li, Yuquan Xie, Rui Shao, Gongwei Chen, Dongmei Jiang, and Liqiang Nie. Optimus-1: Hybrid multimodal memory empowered agents excel in long-horizon tasks. In *Advances in Neural Information Processing Systems*, 2024.

Zaijing Li, Yuquan Xie, Rui Shao, Gongwei Chen, Dongmei Jiang, and Liqiang Nie. Optimus-2: Multimodal minecraft agent with goal-observation-action conditioned policy. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025b.

Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and Andy Zeng. Code as policies: Language model programs for embodied control. In *IEEE International Conference on Robotics and Automation*, 2023.

Shalev Lifshitz, Keiran Paster, Harris Chan, Jimmy Ba, and Sheila McIlraith. Steve-1: A generative model for text-to-behavior in minecraft. In *Advances in neural information processing systems*, volume 36, pages 69900–69929, 2023.

Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling. In *International Conference on Learning Representations*, 2023.

Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and transfer data with rectified flow. In *International Conference on Learning Representations*, 2023.

Zichen Liu, Changyu Chen, Wenjun Li, Penghui Qi, Tianyu Pang, Chao Du, Wee Sun Lee, and Min Lin. Understanding R1-Zero-like training: A critical perspective. *arXiv preprint arXiv:2503.20783*, 2025.

Loïc Magne, Anas Awadalla, Guanzhi Wang, Yinzhen Xu, Joshua Belofsky, Fengyuan Hu, Joohwan Kim, Ludwig Schmidt, Georgia Gkioxari, Jan Kautz, Yisong Yue, Yejin Choi, Yuke Zhu, and Linxi Fan. Nitrogen: An open foundation model for generalist gaming agents. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 21511–21521, 2026.

NVIDIA, Johan Bjorck, Fernando Castaneda, Nikita Cherniadev, Xingye Da, Runyu Ding, Linxi Fan, Yu Fang, Dieter Fox, Fengyuan Hu, et al. GR00T N1: An open foundation model for generalist humanoid robots. *arXiv preprint arXiv:2503.14734*, 2025.

谷歌 DeepMind。Gemini 3.5 Flash 模型卡。https://deepmind.google/models/model-cards/gemini-3-5-flash/，2026a年5月。谷歌 DeepMind。Gemini 3.6 Flash 模型卡。https://deepmind.google/models/model-cards/gemini-3-6-flash/，2026b年7月。William H. Guss、Brandon Houghton、Nicholay Topin、Phillip Wang、Cayden Codel、Manuela Veloso 和 Ruslan Salakhutdinov。MineRL：一个大规模《我的世界》演示数据集。载于国际人工智能联合会议，2019年。Danijar Hafner、Jurgis Pasukonis、Jimmy Ba 和 Timothy Lillicrap。通过世界模型掌握多种控制任务。自然，640(8059):647–653，2025年4月。doi: 10.1038/s41586-025-08744-2。https://doi.org/10.1038/ s41586-025-08744-2。Edward J. Hu、Yelong Shen、Phillip Wallis、Zeyuan Allen-Zhu、Yuanzhi Li、Shean Wang、Lu Wang 和 Weizhu Chen。LoRA：大语言模型的低秩适配。载于国际学习表征会议，2022年。Wenlong Huang、Fei Xia、Ted Xiao、Harris Chan、Jacky Liang、Pete Florence、Andy Zeng、Jonathan Tompson、Igor Mordatch、Yevgen Chebotar 等。内心独白：通过语言模型规划实现具身推理。载于机器人学习会议，2022年。丹尼尔·卡尼曼。《思考，快与慢》。Farrar, Straus and Giroux，2011年。Moo Jin Kim、Karl Pertsch、Siddharth Karamcheti、Ted Xiao、Ashwin Balakrishna、Suraj Nair、Rafael Rafailov、Ethan Foster、Grace Lam、Pannag Sanketi 等。OpenVLA：一个开源视觉-语言-动作模型。arXiv 预印本 arXiv:2406.09246，2024年。Moo Jin Kim、Chelsea Finn 和 Percy Liang。微调视觉-语言-动作模型：优化速度与成功率。arXiv 预印本 arXiv:2502.19645，2025年。Muyao Li、Zihao Wang、Kaichen He、Xiaojian Ma 和 Yitao Liang。JARVIS-VLA：对大规模视觉语言模型进行后训练，以使用键盘和鼠标玩视觉游戏。载于计算语言学协会研究成果：ACL 2025，第17878–17899页，2025a。Zaijing Li、Yuquan Xie、Rui Shao、Gongwei Chen、Dongmei Jiang 和 Liqiang Nie。Optimus-1：混合多模态记忆赋能智能体在长时程任务中表现出色。载于神经信息处理系统进展，2024年。Zaijing Li、Yuquan Xie、Rui Shao、Gongwei Chen、Dongmei Jiang 和 Liqiang Nie。Optimus-2：采用目标-观测-动作条件策略的多模态《我的世界》智能体。载于IEEE/CVF计算机视觉与模式识别会议，2025b。Jacky Liang、Wenlong Huang、Fei Xia、Peng Xu、Karol Hausman、Brian Ichter、Pete Florence 和 Andy Zeng。代码即策略：用于具身控制的语言模型程序。载于IEEE国际机器人与自动化会议，2023年。Shalev Lifshitz、Keiran Paster、Harris Chan、Jimmy Ba 和 Sheila McIlraith。STEVE-1：用于《我的世界》中文本到行为的生成模型。载于神经信息处理系统进展，第36卷，69929，2023年 69900页 – 。Yaron Lipman、Ricky T. Q. Chen、Heli Ben-Hamu、Maximilian Nickel 和 Matt Le。用于生成式建模的流匹配。载于国际学习表征会议，2023年。Xingchao Liu、Chengyue Gong 和 Qiang Liu。笔直而快速地流动：学习使用校正流生成和迁移数据。载于国际学习表征会议，2023年。Zichen Liu、Changyu Chen、Wenjun Li、Penghui Qi、Tianyu Pang、Chao Du、Wee Sun Lee 和 Min Lin。理解类R1-Zero训练：一个批判性视角。arXiv 预印本 arXiv:2503.20783，2025年。Loïc Magne、Anas Awadalla、Guanzhi Wang、Yinzhen Xu、Joshua Belofsky、Fengyuan Hu、Joohwan Kim、Ludwig Schmidt、Georgia Gkioxari、Jan Kautz、Yisong Yue、Yejin Choi、Yuke Zhu 和 Linxi Fan。Nitrogen：面向通用游戏智能体的开放基础模型。载于IEEE/CVF计算机视觉与模式识别会议（CVPR）论文集，第21511–21521页，2026年。英伟达、Johan Bjorck、Fernando Castaneda、Nikita Cherniadev、Xingye Da、Runyu Ding、Linxi Fan、Yu Fang、Dieter Fox、Fengyuan Hu 等。GR00T N1：面向通用人形机器人的开放基础模型。arXiv 预印本 arXiv:2503.14734，2025年。

OpenAI. GPT-4o system card. *arXiv preprint arXiv:2410.21276*, 2024.

OpenAI, Christopher Berner, Greg Brockman, Brooke Chan, Vicki Cheung, Przemysław Dębiak, Christy Dennison, David Farhi, Quirin Fischer, Shariq Hashme, Chris Hesse, Rafal Józefowicz, Scott Gray, Catherine Olsson, Jakub Pachocki, Michael Petrov, Henrique P. d. O. Pinto, Jonathan Raiman, Tim Salimans, Jeremy Schlatter, Jonas Schneider, Szymon Sidor, Ilya Sutskever, Jie Tang, Filip Wolski, and Susan Zhang. Dota 2 with large scale deep reinforcement learning. *arXiv preprint arXiv:1912.06680*, 2019. https://arxiv.org/abs/1912.06680.

Karl Pertsch, Kyle Stachowicz, Brian Ichter, Danny Driess, Suraj Nair, Quan Vuong, Oier Mees, Chelsea Finn, and Sergey Levine. FAST: Efficient action tokenization for vision-language-action models. *arXiv preprint arXiv:2501.09747*, 2025.

Physical Intelligence, Kevin Black, Noah Brown, James Darpinian, Karan Dhabalia, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, et al. $\pi_{0.5}$: A vision-language-action model with open-world generalization. *arXiv preprint arXiv:2504.16054*, 2025.

Qwen Team. Qwen3-Omni technical report. *arXiv preprint arXiv:2509.17765*, 2025.

Qwen Team. Qwen3.5: Towards native multimodal agents. https://qwen.ai/blog?id=qwen3.5, February 2026a.

Qwen Team. Qwen3.6-35B-A3B: Agentic coding power, now open to all. https://qwen.ai/blog?id=qwen3.6-35b-a3b, April 2026b.

Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap, and David Silver. Mastering Atari, Go, chess and shogi by planning with a learned model. *Nature*, 588:604–609, 2020. doi: 10.1038/s41586-020-03051-4.

Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.

Yide Shentu, Philipp Wu, Aravind Rajeswaran, and Pieter Abbeel. From LLMs to actions: Latent codes as bridges in hierarchical robot control. In *IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2024.

Lucy Xiaoyang Shi, Brian Ichter, Michael Equi, Liyiming Ke, Karl Pertsch, Quan Vuong, James Tanner, Anna Walling, Haohuan Wang, Niccolo Fusai, et al. Hi robot: Open-ended instruction following with hierarchical vision-language-action models. In *International Conference on Machine Learning*, 2025.

Mustafa Shukor, Dana Aubakirova, Francesco Capuano, Pepijn Kooijmans, Steven Palma, Adil Zouitine, Michel Aractingi, Caroline Pascal, Martino Russi, Andres Marafioti, Simon Alibert, Matthieu Cord, Thomas Wolf, and Remi Cadene. SmolVLA: A vision-language-action model for affordable and efficient robotics. *arXiv preprint arXiv:2506.01844*, 2025.

David Silver, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. Mastering the game of Go with deep neural networks and tree search. *Nature*, 529:484–489, 2016. doi: 10.1038/nature16961.

David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George van den Driessche, Thore Graepel, and Demis Hassabis. Mastering the game of Go without human knowledge. *Nature*, 550: 354–359, 2017. doi: 10.1038/nature24270.

SIMA Team. Sima 2: A generalist embodied agent for virtual worlds. *arXiv preprint arXiv:2512.04797*, 2025.

Weihao Tan, Wentao Zhang, Xinrun Xu, Haochong Xia, Ziluo Ding, Boyu Li, Bohan Zhou, Junpeng Yue, Jiechuan Jiang, Yewen Li, et al. Cradle: Empowering foundation agents towards general computer control. *arXiv preprint arXiv:2403.03186*, 2024.

Weihao Tan, Xiangyang Li, Yunhao Fang, Heyuan Yao, Shi Yan, Hao Luo, Tenglong Ao, Huihui Li, Hongbin Ren, Bairen Yi, et al. Lumine: An open recipe for building generalist agents in 3D open worlds. *arXiv preprint arXiv:2511.08892*, 2025.

Oriol Vinyals, Igor Babuschkin, Wojciech M. Czarnecki, Michael Mathieu, Andrew Dudzik, Junyoung Chung, David H. Choi, Richard Powell, Timo Ewalds, Petko Georgiev, Junhyuk Oh, Dan Horgan, Manuel Kroiss, Ivo Danihelka, Aja Huang, Laurent Sifre, Trevor Cai, John P. Agapiou, Max Jaderberg, Alexander S. Vezhnevets, Remi Leblond,

OpenAI．GPT-4o 系统卡．*arXiv* 预印本 *arXiv:2410.21276*，2024年。

OpenAI、Christopher Berner、Greg Brockman、Brooke Chan、Vicki Cheung、Przemysław Dębiak、Christy Dennison、David Farhi、Quirin Fischer、Shariq Hashme、Chris Hesse、Rafal Józefowicz、Scott Gray、Catherine Olsson、Jakub Pachocki、Michael Petrov、Henrique P. d. O. Pinto、Jonathan Raiman、Tim Salimans、Jeremy Schlatter、Jonas Schneider、Szymon Sidor、伊利亚·苏茨克维、Jie Tang、Filip Wolski 和 Susan Zhang。基于大规模深度强化学习的 Dota 2。*arXiv* 预印本 *arXiv:1912.06680*，2019年。 https://arxiv.org/abs/1912.06680。

Karl Pertsch、Kyle Stachowicz、Brian Ichter、Danny Driess、Suraj Nair、Quan Vuong、Oier Mees、Chelsea Finn 和 Sergey Levine。FAST：面向视觉-语言-动作模型的高效动作词元化。*arXiv* 预印本 *arXiv:2501.09747*，2025年。

Physical Intelligence, Kevin Black, Noah Brown, James Darpinian, Karan Dhabalia, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, 等. $\pi_{0.5}$：具备开放世界泛化能力的视觉-语言-动作模型。*arXiv* 预印本 *arXiv:2504.16054*，2025年。

通义千问团队。Qwen3-Omni 技术报告。*arXiv* 预印本 *arXiv:2509.17765*，2025年。

通义千问团队。Qw en3.5：迈向原生多模态智能体。https://qwen.ai/blog?id=qwen3.5，202年2月 6a.

通义千问团队。Qwen3.6-35B-A3B：智能体编程能力，现向所有人开放。https://qwen.ai/blog?id=qwen3.6-35b-a3b，2026年4月b。

Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, 德米斯·哈萨比斯, Thore Graepel, Timothy Lillicrap 和 戴维·西尔弗。通过使用学习得到的模型进行规划，精通雅达利、围棋、国际象棋和将棋。自然，588:604–609，2020年。doi: 10.1038/s41586-020-03051-4。

Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y . K. Li, Y. Wu 和 Daya Guo。DeepSeekMath：突破开放语言模型的数学推理极限。*arXiv* 预印本 *arXiv:2402.03300*，2024年。

Yide Shentu, Philipp Wu, Aravind Rajeswaran, and Pieter Abbeel. 从大语言模型到动作：潜在编码作为分层机器人控制中的桥梁。载于*IEEE/RSJ*智能机器人与系统国际会议，2024年。

Lucy Xiaoyang Shi, Brian Ichter, Michael Equi, Liyiming Ke, Karl Pertsch, Quan Vuong, James Tanner, Anna Walling, Haohuan Wang, Niccolo Fusai 等。Hi Robot：利用分层视觉-语言-动作模型实现开放式指令遵循。载于国际机器学习会议，2025年。

Mustafa Shukor、Dana Aubakirova、Francesco Capuano、Pepijn Kooijmans、Steven Palma、Adil Zouitine、Michel Aractingi、Caroline Pascal、Martino Russi、Andres Marafioti、Simon Alibert、Matthieu Cord、Thomas Wolf 和 Re mi Cadene。SmolVLA：面向低成本高效机器人技术的视觉-语言-动作模型。*arXiv* 预印本 *arXiv:2506.01844*，2025年。

戴维·西尔弗、黄士杰、Chris J. Maddison、Arthur Guez、Laurent Sifre、George van den Driessche、Julian Schrittwieser、Ioannis Antonoglou、Veda Panneershelvam、Marc Lanctot、Sander Dieleman、Dominik Grewe、John Nham、Nal Kalchbrenner、伊利亚·苏茨克维、Timothy Lillicrap、Madeleine Leach、Koray Kavukcuoglu、Thore Graepel 和 德米斯·哈萨比斯。利用深度神经网络和树搜索精通围棋。自然，529:484–489，2016年。doi: 10.1038/nature16961。

戴维·西尔弗、Julian Schrittwieser、Karen Simonyan、Ioannis Antonoglou、黄士杰、Arthur Guez、Thomas Hubert、Lucas Baker、Matthew Lai、Adrian Bolton、Yutian Chen、Timothy Lillicrap、樊麾、Laurent Sifre、George van den Driessche、Thore Graepel和德米斯·哈萨比斯。无需人类知识即可精通围棋。自然，550：354–359，2017年。doi: 10.1038/nature24270.

SIMA团队．Sima 2：面向虚拟世界的通用具身智能体．*arXiv*预印本 *arXiv:2512.04797*，2025年．

Weihao Tan、Wentao Zhang、Xinrun Xu、Haochong Xia、Ziluo Ding、Boyu Li、Bohan Zhou、Junpeng Yue、Jiechuan Jiang、Yewen Li等。Cradle：赋能基础智能体，实现通用计算机控制。*arXiv*预印本 *arXiv:2403.03186*，2024年。

Weihao Tan, Xiangyang Li, Yunhao Fang, Heyuan Yao, Shi Yan, Hao Luo, Tenglong Ao, Huihui Li, Hongbin Ren, Bairen Yi 等。Lumine：在三维开放世界中建造通用智能体的开放方案。*arXiv* 预印本 *arXiv:2511.08892*，2025年。

Oriol Vinyals, Igor Babuschkin, Wojciech M. Czarnecki, Michael Mathieu, Andrew Dudzik, JunyoungChung, David H. Choi, Richard Powell, Timo Ewalds, Petko Georgiev, Junhyuk Oh, Dan Horgan, Manuel Kroiss, Ivo Danihelka, 黄士杰, Laurent Sifre, Trevor Cai, John P. Agapiou, Max Jaderberg, Alexander S. Vezhnevets, Remi Leblond,

Tobias Pohlen, Valentin Dalibard, David Budden, Yury Sulsky, James Molloy, Tom L. Paine, Caglar Gulcehre, Ziyu Wang, Tobias Pfaff, Yuhuai Wu, Roman Ring, Dani Yogatama, Dario Wunsch, Katrina McKinney, Oliver Smith, Tom Schaul, Timothy Lillicrap, Koray Kavukcuoglu, Demis Hassabis, Chris Apps, and David Silver. Grandmaster level in StarCraft II using multi-agent reinforcement learning. *Nature*, 575:350–354, 2019.

Wan Team. Wan: Open and advanced large-scale video generative models. *arXiv preprint arXiv:2503.20314*, 2025.

Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023a.

Yiran Wang, Xingyilang Yin, Junfu Pu, Guangzhi Wang, Kaifeng Li, Mingyu Ouyang, Huiqiang Sun, Lingen Li, Cheng Cheng, Wangbo Yu, Honghao Chen, Xiaodong Cun, Chi-Man Pun, Zhiguo Cao, and Ying Shan. GameHorizon Suite: Multi-horizon data and evaluation in gameplay. *arXiv preprint arXiv:2609.25001*, 2026. https://arxiv.org/abs/2609.25001.

Zihao Wang, Shaofei Cai, Guanzhou Chen, Anji Liu, Xiaojian Ma, and Yitao Liang. Describe, explain, plan and select: Interactive planning with large language models enables open-world multi-task agents. *arXiv preprint arXiv:2302.01560*, 2023b.

Zihao Wang, Shaofei Cai, Anji Liu, Yonggang Jin, Jinbing Hou, Bowei Zhang, Haowei Lin, Zhaofeng He, Zilong Zheng, Yaodong Yang, Xiaojian Ma, and Yitao Liang. JARVIS-1: Open-world multi-task agents with memory-augmented multimodal language models. *arXiv preprint arXiv:2311.05997*, 2023c.

Zihao Wang, Shaofei Cai, Zhancun Mu, Haowei Lin, Ceyao Zhang, Xuejie Liu, Qing Li, Anji Liu, Xiaojian Ma, and Yitao Liang. OmniJARVIS: Unified vision-language-action tokenization enables open-world instruction following agents. In *Advances in Neural Information Processing Systems*, 2024.

Zihao Wang, Xujing Li, Yining Ye, Junjie Fang, Haoming Wang, Longxiang Liu, Shihao Liang, Junting Lu, Zhiyong Wu, Jiazhan Feng, et al. Game-tars: Pretrained foundation models for scalable generalist multimodal game agents. *arXiv preprint arXiv:2510.23691*, 2025.

Shenhao Yan, Ge Wang, Qi Liu, Weilin Meng, Jiahao Yang, Chengsi Yao, Fan Feng, Xiaoguang Ma, Yiming Zhao, and Yatong Han. Acting while understanding: Asynchronous semantic-action decoupling for real-time vision-language-action models. *arXiv preprint arXiv:2606.15285*, 2026.

Haoqi Yuan, Chi Zhang, Hongcheng Wang, Feiyang Xie, Penglin Cai, Hao Dong, and Zongqing Lu. Skill reinforcement learning and planning for open-world long-horizon tasks. *arXiv preprint arXiv:2303.16563*, 2023.

Yuguang Yue, Irakli Salia, Samuel Hunt, Chris Green, Wenzhe Shi, and Jonathan J Hunt. Scaling behavior cloning improves causal reasoning: An open model for real-time video game playing. *arXiv preprint arXiv:2601.04575*, 2026.

Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language image pre-training. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 11975–11986, 2023.

Jianke Zhang, Yanjiang Guo, Xiaoyu Chen, Yen-Jen Wang, Yucheng Hu, Chengming Shi, and Jianyu Chen. HiRT: Enhancing robotic control with hierarchical robot transformers. In *Conference on Robot Learning*, 2024.

Tony Z. Zhao, Vikash Kumar, Sergey Levine, and Chelsea Finn. Learning fine-grained bimanual manipulation with low-cost hardware. In *Robotics: Science and Systems*, 2023.

Xinyue Zheng, Haowei Lin, Kaichen He, Zihao Wang, Qiang Fu, Haobo Fu, Zilong Zheng, and Yitao Liang. MCU: An evaluation framework for open-ended game agents. In *International conference on machine learning*, pages 78221–78259. PMLR, 2025.

Enshen Zhou, Yiran Qin, Zhenfei Yin, Yuzhou Huang, Ruimao Zhang, Lu Sheng, Yu Qiao, and Jing Shao. MineDreamer: Learning to follow instructions via chain-of-imagination for simulated-world control. *arXiv preprint arXiv:2403.12037*, 2024.

Tobias Pohlen、Valentin Dalibard、David Budden、Yury Sulsky、James Molloy、Tom L. Paine、Caglar Gulcehre、Ziyu Wang、Tobias Pfaff、Yuhuai Wu、Roman Ring、Dani Yogatama、Dario Wunsch、Katrina McKinney、Oliver Smith、Tom Schaul、Timothy Lillicrap、Koray Kavukcuoglu、德米斯·哈萨比斯、Chris Apps和戴维·西尔弗。利用多智能体强化学习在星际争霸II中达到宗师级水平。自然，575:350–354，2019年。

Wan团队 . 万相：开放且先进的大规模视频生成模型 . *arXiv*预印本 *arXiv:2503* .*20314*，2025年.

Guanzhi Wang、Yuqi Xie、Yunfan Jiang、Ajay Mandlekar、Chaowei Xiao、Yuke Zhu、Linxi Fan和Anima Anandkumar。Voyager：基于大语言模型的开放式具身智能体。*arXiv*预印本 *arXiv:2305.16291*，2023年a。

Yiran Wang、Xingyilang Yin、Junfu Pu、Guangzhi Wang、Kaifeng Li、Mingyu Ouyang、Huiqiang Sun、Lingen Li、Cheng Cheng、Wangbo Yu、Honghao Chen、Xiaodong Cun、Chi-Man Pun、Zhiguo Cao 和 Ying Shan。GameHorizon Suite：游戏过程中的多时间跨度数据与评估。*arXiv* 预印本 *arXiv:2609.25001*，2026年。https://arxiv.org/abs/2609.25001。

Zihao Wang、Shaofei Cai、Guanzhou Chen、Anji Liu、Xiaojian Ma 和 Yitao Liang。描述、解释、规划与选择：基于大语言模型的交互式规划赋能开放世界多任务智能体。*arXiv* 预印本 *arXiv:2302.01560*，2023b。

Zihao Wang、Shaofei Cai、Anji Liu、Yonggang Jin、Jinbing Hou、Bowei Zhang、Haowei Lin、Zhaofeng He、Zilong Zheng、Yaodong Yang、Xiaojian Ma 和 Yitao Liang。JARVIS-1：基于记忆增强多模态语言模型的开放世界多任务智能体。*arXiv* 预印本 *arXiv:2311.05997*，2023c。

Zihao Wang、Shaofei Cai、Zhancun Mu、Haowei Lin、Ceyao Zhang、Xuejie Liu、Qing Li、Anji Liu、Xiaojian Ma 和 Yitao Liang。OmniJARVIS：统一视觉-语言-动作词元化赋能开放世界指令遵循智能体。载于神经信息处理系统进展，2024年。

Zihao Wang、Xujing Li、Yining Ye、Junjie Fang、Haoming Wang、Longxiang Liu、Shihao Liang、Junting Lu、Zhiyong Wu、Jiazhan Feng 等。Game-tars：面向可扩展通用多模态游戏智能体的预训练基础模型。*arXiv* 预印本 *arXiv:2510.23691*，2025年。

Shenhao Yan、Ge Wang、Qi Liu、Weilin Meng、Jiahao Yang、Chengsi Yao、Fan Feng、Xiaoguang Ma、Yiming Zhao 和 Yatong Han。Acting While Understanding：面向实时视觉-语言-动作模型的异步语义-动作解耦。*arXiv* 预印本 *arXiv:2606.15285*，2026年。

Haoqi Yuan、Chi Zhang、Hongcheng Wang、Feiyang Xie、Penglin Cai、Hao Dong 和 Zongqing Lu。面向开放世界长时程任务的技能强化学习与规划。*arXiv*预印本 *arXiv:2303.16563*，2023年。

Yuguang Yue、Irakli Salia、Samuel Hunt、Chris Green、Wenzhe Shi 和 Jonathan J Hunt。扩大行为克隆规模可提升因果推理能力：用于实时电子游戏游玩的开放模型。*arXiv*预印本 *arXiv:2601.04575*，2026年。

Xiaohua Zhai、Basil Mustafa、Alexander Kolesnikov 和 Lucas Beyer。用于语言-图像预训练的Sigmoid损失。载于*IEEE/CVF*国际计算机视觉会议论文集，第11975–11986页，2023年。

Jianke Zhang、Yanjiang Guo、Xiaoyu Chen、Yen-Jen Wang、Yucheng Hu、Chengming Shi 和 Jianyu Chen。HiRT：利用分层机器人Transformer增强机器人控制。载于机器人学习会议，2024年。Tony Z. Zhao、Vikash Kumar、Sergey Levine 和 Chelsea Finn。利用低成本硬件学习细粒度双臂操作。载于机器人学：科学与系统，2023年。

Xinyue Zheng、Haowei Lin、Kaichen He、Zihao Wang、Qiang Fu、Haobo Fu、Zilong Zheng 和 Yitao Liang . MCU：面向开放式游戏智能体的评估框架。载于国际机器学习会议，第 78221–78259 页。PMLR，2025年。

Enshen Zhou、Yiran Qin、Zhenfei Yin、Yuzhou Huang、Ruimao Zhang、Lu Sheng、Yu Qiao 和 Jing Shao . MineDreamer：通过想象链学习遵循指令以实现模拟世界控制。*arXiv* 预印本 *arXiv:2403 12037*，2024年。。

Appendix

A Task Planner Details

A.1 Decomposition data

We use the Minecraft decomposition corpus of GameHorizon Suite (Wang et al., 2026); its construction, annotation procedure, and statistics are described there. We split it by complete video clip, yielding 57,901 training examples from 1,119 clips and 2,200 held-out examples from 50 disjoint clips; the held-out set is the SSF1 evaluation set of Section 4. After adding the second description style and the paraphrased instructions and downsampling thematic examples without replacement, the SFT training set contains 40,068 examples: 10,442 causal, 2,914 original-segment, and 26,712 thematic.

A.2 Evaluation protocols

Soft Semantic F1. The model receives one observation frame and a high-level instruction without the reference step count. Predicted and reference steps are encoded with `all-mpnet-base-v2`, and pairs with cosine similarity below 0.65 are excluded before alignment. Causal examples use order-preserving one-to-one alignment, whereas thematic examples use unrestricted Hungarian assignment. For n reference steps, m predicted steps, and matched pairs $\mathcal{M}$,

$$\mathrm{SSF1} = \frac{2\sum_{(i,j)\in\mathcal{M}} \cos(e_i, \hat{e}_j)}{n+m}. \tag{16}$$

We report the mean over the 2,200 held-out examples; unparseable outputs receive zero.

Task2. We use the high-level-to-atomic decomposition task (Task2) of GameHorizon Suite (Wang et al., 2026) with a generate-and-align protocol on 250 Minecraft questions. The model receives one observation frame, a high-level instruction, and the target step count, and generates a numbered plan. Each question provides four permutations of the same canonical step set. Generated steps are aligned with the canonical steps using the similarity in Equation (4), the aligned sequence is compared with each candidate by Kendall correlation, and the unique highest-scoring candidate is selected. The Unparsed rate is the percentage of questions without a valid decision, including ties; these count as incorrect.

A.3 Training hyperparameters

Table 6 Planner training hyperparameters.

	SFT	GRPO
Initialization	Qwen3.5-9B	SFT checkpoint
Trainable modules	language model, connector	language model, connector
Training data	40,068 examples, 1 epoch	2,207 examples
Peak learning rate	2×10^{-5}	1×10^{-6}
Schedule	cosine, 0.5% warmup	cosine, 3% warmup
Weight decay	0.01	–
Batch	4 per GPU × 8 GPUs	32 generations per step
Generations per input	–	8
KL coefficient β	–	0.04
Max. sequence / generation length	8,192	512
Gradient clipping	–	1.0
Parallelism	ZeRO-3	ZeRO-3, colocated inference (TP=2)

The vision encoder is frozen in both stages. Following Liu et al. (2025), GRPO uses mean-centered group rewards without dividing by the within-group standard deviation: when all candidates in a group score

附录

A TaskPlanner 详细说明

A.1 分解数据

我们使用 GameHorizon Suite 的《我的世界》分解语料库（Wang 等人，2026年）；其构建方法、标注流程和统计信息均见该文献。我们按完整视频片段进行划分，从 1,119 个片段中获得 57,901 个训练样本，从另外 50 个互不重叠的片段中获得 2,200 个留出样本；该留出集即为第 4 节中的 SSF1 评估集。在添加第二种描述风格和改写指令，并对主题型样本进行无放回下采样后，监督微调训练集包含 40,068 个样本：10,442 个因果型样本、2,914 个原始片段型样本和 26,712 个主题型样本。

A.2 评估协议

软语义 *F1*。模型接收一个观测帧和一条高层指令，不提供参考步骤数。预测步骤与参考步骤使用 all-mpnet-base-v2 进行编码，余弦相似度低于 0.65 的配对在对齐前被排除。因果型样本采用保序一对一对齐，而主题型样本采用无约束匈牙利指派。对于 n 个参考步骤、 m 个预测步骤以及匹配对 $\mathcal{M}$,

$$\mathrm{SSF1} = \frac{2\sum_{(i,j)\in\mathcal{M}} \cos(e_i, \hat{e}_j)}{n+m}. \tag{16}$$

我们报告在 2,200 个留出样本上的平均值；无法解析的输出计为零分。

Task2。我们采用 GameHorizon Suite（Wang 等人，2026年）中的高层到原子级分解任务（Task2），在 250 道《我的世界》问题上使用生成并对齐协议。模型接收一个观测帧、一条高层指令以及目标步骤数，并生成编号计划。每道问题提供同一标准步骤集的四种排列。使用公式（4）中的相似度将生成步骤与标准步骤对齐，通过肯德尔相关系数将对齐序列与各候选项进行比较，并选出唯一得分最高的候选项。未解析率是未能得出有效判定的问题所占的百分比，包括出现平分的情况；这些情况均计为错误。

A.3 训练超参数

表 6 规划器训练超参数。

	SFT	GRPO
初始化	Qwen3.5-9B	监督微调检查点
可训练模块	语言模型、连接器	语言模型、连接器
训练数据	40,068 个样本，1 个训练轮次	2,207 个样本
峰值学习率	2×10^{-5}	1×10^{-6}
学习率调度	余弦调度，0.5% 预热	余弦调度，3% 预热
权重衰减	0.01	–
批次	每个图形处理器 4 个 × 8个图形处理器	每步生成次数为32
每个输入的生成次数	–	8
KL系数 β	–	0.04
最大序列/生成长度	8,192	512
梯度裁剪	–	1.0
并行度	ZeRO-3	ZeRO-3，同置推理（TP=2）

两个阶段均冻结视觉编码器。遵循Liu等人（2025年）的做法，组相对策略优化使用均值中心化组奖励，不除以组内标准差：当组内所有候选的得分

similarly, a very small standard deviation would amplify negligible reward differences into large advantages, whereas unnormalized rewards preserve the intended relative scale of the count and order terms.

A.4 Dependency annotations and reward weights

The dependency annotations of GameHorizon Suite (Wang et al., 2026) show that causal plans are much more constrained than thematic ones, and that a large fraction of thematic plans are essentially unordered, so a single canonical-order reward would over-constrain the thematic majority.

To set the weights of Equation (10), we simulate the value of each ordering signal on the annotated graphs. A uniformly random permutation reaches 19.8% on the Task2 ordering criterion. Satisfying all annotated dependencies adds 30.8 points, whereas additionally following the recorded order on the remaining free pairs adds only 4.6 points. The implied contributions of dependency and conventional ordering are approximately 87:13, which motivates the 0.8/0.2 split between S_{dep} and S_{conv}.

A.5 GRPO training dynamics

Table 7 GRPO training statistics over the first 450 steps (means within each window).

Steps	Total reward	Count score	Order score	KL	Grad. norm
1–50	0.313	0.630	0.178	0.003	1.669
51–150	0.331	0.664	0.188	0.021	1.619
151–250	0.356	0.694	0.211	0.050	1.408
251–350	**0.379**	0.702	**0.240**	0.044	1.496
351–450	0.369	**0.704**	0.226	0.040	1.150

The count score rises from 0.630 to 0.704, and the order score improves from 0.178 to 0.226–0.240 before plateauing. KL to the SFT policy stays between 0.02 and 0.05 after warmup, and the gradient norm decreases without signs of policy collapse. Generated plans average 73–86 tokens and never reach the 512-token limit, so count errors are not truncation artifacts.

A.6 Why naive SFT generalizes poorly

Our first SFT run balanced categories by duplicating minority examples. Its training loss decreased steadily, but on an external diagnostic built from VPT-20h (539 four-way questions over 498 clips that share no recordings with our training data; chance is 25%), accuracy fell to 39.5%, below the backbone's 51.2%. Grouping questions by how closely the gold steps match the phrasing of the training annotations shows that the SFT model beats the backbone on the most familiar quartile (76.1% vs. 70.1%) but falls below chance on the least familiar one (17.0% vs. 35.6%). The model thus learns a narrow decomposition skill tied to annotation phrasing, which a clip-level holdout from the training distribution did not expose. These findings motivated our SFT recipe (Section A.1): we balance categories by downsampling thematic examples instead of duplicating minority ones, describe each atomic target in two styles, and paraphrase high-level instructions.

B Frame Controller Details

B.1 Training data

The controller is trained on the VPT-100h gameplay of GameHorizon Suite (Wang et al., 2026), about 100 hours of first-person Minecraft play by human contractors (Baker et al., 2022) with per-frame keyboard, mouse, and camera actions and atomic-instruction annotations; we use its 1,132 training clips and 35 validation clips at 20 fps and refer to it for data collection and annotation details. Frames are resized from 640×360 to 256×256 without cropping to preserve the full field of view, and GUI frames (inventory and crafting screens) are excluded. We split each clip into *clean segments* in which the instruction is constant and no GUI frame

同样，极小的标准差会将微不足道的奖励差异放大为很大的优势值，而未归一化奖励则保留了计数项与顺序项预期的相对尺度。

A.4 依赖关系标注与奖励权重

GameHorizon Suite 的依赖关系标注（Wang 等人，2026年）表明，因果计划受到的约束远多于主题计划，而且相当大比例的主题计划基本上是无序的，因此，单一的规范顺序奖励会对占多数的主题计划施加过度约束。

为设定公式 (10) 的权重，我们在标注图上模拟各排序信号的取值。均匀随机排列在 Task2 排序标准下的得分为 19.8%。满足所有已标注的依赖关系可使得分提高 30.8 个百分点，而在此基础上，对其余自由配对也遵循记录顺序，仅能再提高 4.6 个百分点。由此推算，依赖关系与惯常顺序的贡献之比约为 87:13，因此我们将 S_{dep} 和 S_{conv}的权重分别设为 0.8 和 0.2。

A.5 组相对策略优化训练动态

表 7 前 450 步的 GRPO 训练统计（各窗口内的均值）。

步数	总奖励	计数得分	顺序得分	KL	梯度范数
1–50	0.313	0.630	0.178	0.003	1.669
51–150	0.331	0.664	0.188	0.021	1.619
151–250	0.356	0.694	0.211	0.050	1.408
251–350	**0.379**	0.702	**0.240**	0.044	1.496
351–450	0.369	**0.704**	0.226	0.040	1.150

计数得分从 0.630 上升至 0.704，顺序得分从 0.178 提升至 0.226–0.240，随后趋于平稳。预热后，相对于监督微调策略的 KL 散度保持在 0.02 至 0.05 之间，梯度范数下降，且没有出现策略崩溃的迹象。生成计划的平均长度为 73–86 个词元，从未达到 512 个词元的上限，因此计数错误并非截断伪影。

A.6 为什么朴素监督微调的泛化能力较差

我们在首次监督微调中，通过复制少数类样本来平衡类别。其训练损失稳步下降，但在基于 VPT-20h 构建的外部诊断测试中（涵盖 498 个片段的 539 道四选一问题，这些片段与我们的训练数据不共享任何原始录像；随机猜测的准确率为 25%），准确率降至 39.5%，低于骨干网络的 51.2%。按照标准答案步骤与训练标注措辞的匹配程度对问题进行分组后发现，监督微调模型在最熟悉的四分位组中优于骨干网络（76.1% 对 70.1%），但在最不熟悉的四分位组中，其准确率甚至低于随机猜测水平（17.0% 对 35.6%）。因此，该模型学到的是一种与标注措辞绑定的狭隘分解能力，而从训练分布中划分出的片段级留出集未能揭示这一问题。这些发现促使我们制定了相应的监督微调方案（第 A.1 节）：通过对主题型样本进行下采样而非复制少数类样本来平衡类别，以两种风格描述每个原子目标，并对高层指令进行改写。

B 帧控制器详情

B.1 训练数据

控制器在 GameHorizon Suite 的 VPT-100h 游戏过程数据上进行训练（Wang 等人，2026年），该数据包含约 100 小时由受雇人员以第一人称视角游玩《我的世界》的游戏过程（Baker 等人，2022年），并附有逐帧键盘、鼠标和相机动作以及原子指令标注；我们使用其中的 1,132 个训练片段和 35 个验证片段，帧率为 20 fps，数据采集和标注的详细信息请参见该数据集。帧尺寸从 640×360 调整为256×256，不进行裁剪，以保留完整视野，并排除图形用户界面帧（物品栏和合成界面）。我们将每个片段划分为干净片段，其中指令保持不变，且没有图形用户界面帧

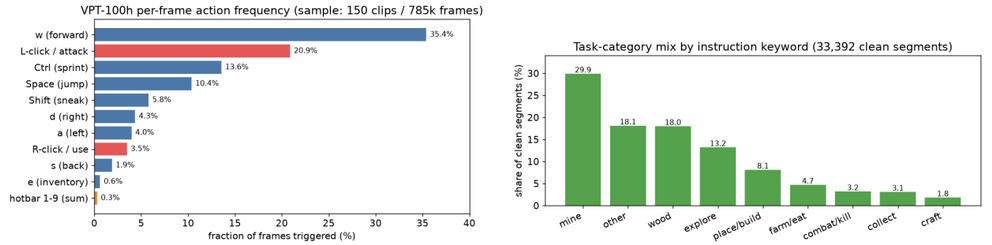

Figure 4 Statistics of the controller training data. **Left:** per-frame action frequency on a sample of 150 clips (785K frames). **Right:** task category mix of clean segments, grouped by instruction keyword.

occurs, giving 33,392 segments; training windows are tiled inside a single segment, which removes the roughly 16% of windows that would otherwise straddle an instruction boundary.

The data is strongly long-tailed (Figure 4). Moving forward occurs in 35.4% of frames and attacking in 20.9%, whereas using or placing items occurs in 3.5% and hotbar selection in only 0.3%; 54.7% of frames contain no key press. Combat accounts for only 3.2% of clean segments. We therefore oversample segments by category and by rare actions: combat segments by a factor of 6, placement and building by 3, and segments containing a use action by 1.5.

B.2 Action space

The keyboard head predicts up to four simultaneous keys from 26 classes: no key, W, A, S, D, space (jump), left and right shift (sneak), left control (sprint), E (inventory), Q (drop), F (swap hands), Z, the four arrow keys, and hotbar slots 1–9. The mouse-button head predicts up to two simultaneous buttons from four classes (none, attack, use, middle). Both are decoded autoregressively from the action-output token. The flow-matching head predicts camera motion as yaw and pitch increments in degrees, converted from recorded pixel displacement at 0.15° per pixel and normalized by per-axis standard deviations estimated from the data (about 5° for yaw and 2° for pitch). Camera targets are masked on GUI frames.

B.3 Architecture and optimization

Auxiliary losses. During joint alignment (Stage 3), the adapted planner is additionally trained to describe the actions taken in the current window as a short sequence of action phrases, with a cross-entropy loss (weight 1.5) that up-weights rare action categories such as camera, interaction, and hotbar actions. The reported controller was additionally trained with an auxiliary video-prediction loss (weight 0.1): a branch initialized from Wan2.1-T2V-1.3B (Wan Team, 2025) predicted 8 future frames from the controller's hidden states and was discarded at inference, so it does not affect the policy's architecture or inference cost. In our experiments, removing this loss had little effect on performance, and we do not consider it part of the method.

B.4 Flow-matching camera head

We compare the flow-matching camera head with a discrete head that classifies camera motion into bins and predicts their expected value, following Yue et al. (2026). The two controllers are identical except for the camera head, and both are trained with the same data and schedule and evaluated on the held-out VPT-100h validation split; results are in Table 5.

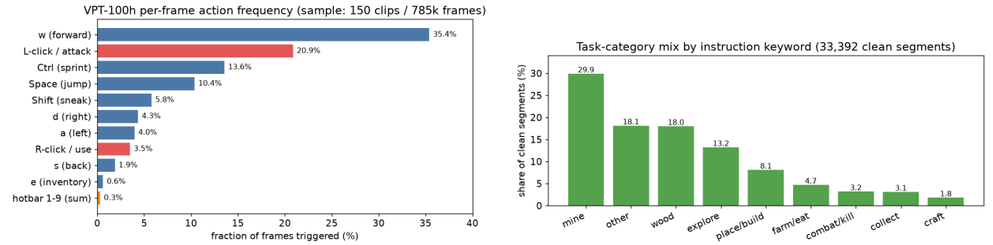

图 4 控制器训练数据的统计信息。**左：** 150 个视频片段样本（78.5 万帧）的逐帧动作频率。**右：** 按指令关键词分组的干净片段任务类别构成。

出现，共得到 33,392 个片段；训练窗口在单个片段内平铺，从而剔除了原本会跨越指令边界的约 16% 的窗口。

数据呈现显著的长尾分布（图 4）。向前移动出现在 35.4% 的帧中，攻击出现在 20.9% 的帧中，而使用或放置物品仅出现在 3.5% 的帧中，快捷栏选择仅出现在 0.3% 的帧中；54.7% 的帧不包含任何按键操作。战斗仅占干净片段的 3.2%。因此，我们按类别和稀有动作对片段进行过采样：战斗片段的采样倍数为 6，放置和建造片段为 3，包含使用动作的片段为 1.5。

B.2 动作空间

键盘预测头从 26 个类别中预测最多四个同时按下的按键：无按键、W、A、S、D、空格键（跳跃）、左右 Shift 键（潜行）、左 Control 键（疾跑）、E（物品栏）、Q（丢弃）、F（切换双手）、Z、四个方向键，以及快捷栏槽位 1–9。鼠标按键预测头从四个类别（无、攻击、使用、中键）中预测最多两个同时按下的按钮。两者均从动作输出词元进行自回归解码。流匹配头以度为单位预测偏航角和俯仰角增量，以表示相机运动；这些增量由记录的像素位移按每像素 0.15° 换算得到，并使用从数据中估计的各轴标准差进行归一化（偏航角约为 5°，俯仰角约为 2°）。图形用户界面帧上的视角目标值会被掩蔽。

B.3 架构与优化

辅助损失。在联合对齐（阶段3）期间，适配后的规划器还接受额外训练，将当前窗口内执行的动作描述为简短的动作短语序列，并采用交叉熵损失（权重为1.5），提高摄像机、交互和快捷栏动作等稀有动作类别的权重。所报告的控制器还使用辅助视频预测损失（权重为0.1）进行了额外训练：一个由Wan2.1-T2V-1.3B（Wan团队，2025年）初始化的分支根据控制器的隐藏状态预测未来8帧，并在推理时被丢弃，因此不会影响策略的架构或推理成本。在我们的实验中，移除该损失对性能几乎没有影响，因此我们不将其视为本方法的一部分。

B.4 流匹配相机预测头

我们遵循 Yue等人（2026年）的方法，将流匹配相机预测头与一种将相机运动划分为分箱并预测其期望值的离散预测头进行比较。两个控制器除相机预测头外完全相同，均使用相同的数据和学习率调度进行训练，并在留出的 VPT-100h 验证集划分上进行评估；结果见表 5。

Table 8 Frame controller hyperparameters.

Backbone	causal Transformer, 30 layers, width 1024, 16 heads
Initialization	pretrained policy of Yue et al. (2026), layers duplicated to 30 with near-identity initialization
Context	$T = 21$ frames ($\approx$1 s), attention window 21 frames
Visual encoder	frozen SigLIP, 256×256 input; 256 patches compressed to 1 token per frame by a Perceiver resampler
Goal condition	hidden state at the goal token, MLP projection
Stage-2 goal encoder	frozen Qwen3.5-9B, current frame and the instruction
Stage-3 planner	GRPO planner (Qwen3.5-9B) with LoRA rank 64, $\alpha = 128$, learning rate 5×10^{-6}
Keyboard / button decoder	autoregressive, 26 / 4 classes
Mouse head	flow matching, hidden size 512, 10 Euler steps
Loss weights	$\lambda_{\text{mouse}} = 1.0$
Optimizer	AdamW, $\beta = (0.9, 0.999)$, weight decay 10^{-4}
Learning rate	5×10^{-5}, 500 warmup steps
Batch	2 per GPU $\times$ 16 GPUs, bf16 mixed precision
Gradient clipping	1.0

C Closed-Loop Evaluation Details

C.1 Task configurations

All tasks run in MineStudio (Cai et al., 2024a) with 256×256 observations for TaskFrame, a 400-step limit, and success defined as the first positive reward from the corresponding event. *Mine* tasks equip the agent with the required pickaxe or axe and fill a hollow $7 \times 5 \times 7$ shell of the target block around it, so a target is visible and reachable from any direction; the instruction is, for example, "Mine stone blocks using a pickaxe." *Kill* tasks equip an iron sword and spawn four target animals 2–6 blocks in front of the agent, with the instruction "Attack and kill the cow with a sword."; the zombie task is run at night with night vision. Task configurations are frozen and shared across all compared models.

C.2 Baseline settings

JARVIS-VLA-7B is served with vLLM and decoded with its own tokenized action space and μ-law camera discretization at 640×360, using the sampling temperature of its official rollout (0.9). *UI-TARS-1.5-7B*, which appears only in the qualitative comparison of Section C.4, receives 640×360 observations with its native interleaved screenshot-and-action history, and each decision is held for 20 game steps because the model emits GUI events at well below 1 Hz. Its prompt describes the effect of each action without suggesting a strategy, and we sample with temperature 0.7, because greedy decoding on near-identical consecutive frames repeats the same action indefinitely. Its GUI events (click, drag, hotkey) are mapped to game actions by a fixed adapter.

C.3 Per-task mining results

Table 9 reports the per-task mining results summarized by the Mine column of Table 1. TaskFrame and JARVIS-VLA-7B succeed in almost every episode, so these tasks do not separate the two models.

C.4 Qualitative comparison with baselines

Figure 6 compares rollouts of TaskFrame with those of the baselines on the same scenes; green and red bars mark successful and failed episodes. We additionally include the unified multimodal models BAGEL (Deng

表8 帧控制器超参数。

骨干网络	因果Transformer，30层，宽度1024，16个注意力头
初始化	预训练策略，来自Yue等人（2026年），通过复制将层数扩展至30层，并采用近恒等初始化
上下文	$T = 21$ 帧（≈1秒），注意力窗口为21帧
视觉编码器	冻结的SigLIP，256×256 输入；256个图像块被压缩由Perceiver重采样器压缩为每帧1个词元
目标条件	目标词元处的隐藏状态，多层感知机投影
阶段2目标编码器	冻结的Qwen3.5-9B，当前帧和指令
阶段3规划器	LoRA秩为64的GRPO规划器（Qwen3.5-9B），$\alpha = 128$，学习率 5×10^{-6}
键盘/按钮解码器	自回归，26 / 4 个类别
鼠标预测头	流匹配，隐藏层维度 512，10 个欧拉步
损失权重	$\lambda_{\text{mouse}} = 1.0$
优化器	AdamW，$\beta = (0.9, 0.999)$，权重衰减 10^{-4}
学习率	5×10^{-5}，500 预热步数
批次	每个图形处理器 2，$\times$ 16 个图形处理器，bf16混合精度
梯度裁剪	1.0

C 闭环评估详情

C.1 任务配置

所有任务均在 MineStudio（Cai 等人，2024a）中运行，为 TaskFrame 提供 256×256 观测，步数上限为 400 步，并将相应事件首次产生正奖励定义为成功。开采任务为智能体配备所需的镐或斧，并在其周围用目标方块构建一个空心的 $7 \times 5 \times 7$ 外壳，使其从任何方向都能看到并触及目标；指令例如“使用镐开采石头方块”。击杀任务为智能体配备铁剑，并在其前方 2–6 个方块处生成四只目标动物，指令为“用剑攻击并杀死牛。”；僵尸任务在夜间运行，并启用夜视。任务配置固定不变，所有参与比较的模型均使用相同配置。

C.2 基线设置

JARVIS-VLA-7B 通过 vLLM 提供服务，使用其自身的词元化动作空间进行解码，并采用 μ 律进行摄像机离散化，设置为 640×360，采样温度沿用其官方轨迹展开中的设置（0.9）。*UI-TARS-1.5-7B* 仅出现在第 C.4 节的定性比较中，它接收 640×360 观测，并使用其原生的交错的截图与动作历史；由于该模型输出图形用户界面事件的频率远低于 1 Hz，每次决策均持续执行 20 个游戏步。其提示词描述了每个动作的效果，但不提供策略建议；我们采用 0.7 的采样温度，因为对几乎相同的连续帧进行贪心解码会导致同一动作无限重复。其图形用户界面事件（点击、拖拽、快捷键）通过固定适配器映射为游戏动作。

C.3 各任务的采矿结果

表 9 报告了表 1中采矿列所汇总的各任务采矿结果。TaskFrame 和 JARVIS-VLA-7B 在几乎所有回合中都取得了成功，因此这些任务无法区分两个模型的表现。

C.4 与基线的定性比较

图 6 比较了 TaskFrame 与各基线在相同场景中的轨迹展开；绿色和红色条分别标示成功和失败的回合。我们还纳入了统一多模态模型 BAGEL（Deng

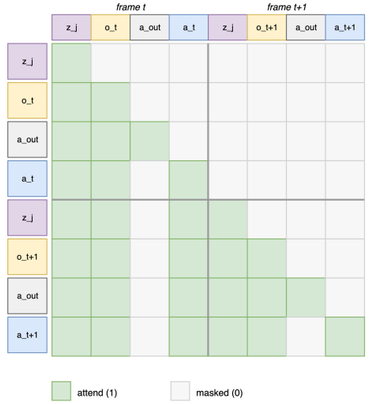

Figure 5 Causal attention mask of the frame controller, adopted from Yue et al. (2026) and shown for two consecutive frames. Rows are query tokens and columns are key tokens; green cells are attended and gray cells are masked. Each frame contributes the goal condition z_j, the visual token o_t, the action-output token a^{out}, and the ground-truth action a_t. No token attends to a^{out}, and a^{out} at frame t attends to previous ground-truth actions but not to the same-frame action a_t.

Table 9 Per-task closed-loop success rate (%) on mining tasks over 3 evaluation seeds (6 episodes per seed), with binomial estimates of the standard deviation across seeds.

Model	Mine Blocks				
	Stone	Iron ore	Coal ore	Oak log	Avg.
VPT-BC (Baker et al., 2022)	67 ±19.6	17 ±16.8	17 ±16.8	50 ±20.4	38 ±9.2
VPT-RL (Baker et al., 2022)	50 ±20.4	0 ±10.5	17 ±16.8	50 ±20.4	29 ±8.8
STEVE-1 (Lifshitz et al., 2023)	83 ±16.8	33 ±19.6	50 ±20.4	83 ±16.8	62 ±9.2
JARVIS-VLA-7B (Li et al., 2025a)	**100** ±10.5	**100** ±10.5	**100** ±10.5	**100** ±10.5	**100** ±5.3
TaskFrame (Ours)	**100** ±10.5	**100** ±10.5	83 ±16.8	**100** ±10.5	96 ±6.2

et al., 2025) and SenseNova-U1 (Diao et al., 2026), run zero-shot through the same action adapter as UI-TARS-1.5-7B. On kill sheep, only TaskFrame succeeds: JARVIS-VLA-7B wanders without engaging the sheep, BAGEL turns the camera toward the sky and does not recover, UI-TARS-1.5-7B turns away from the targets, and SenseNova-U1 approaches a sheep but does not kill it. Mining oak logs is easier: TaskFrame, JARVIS-VLA-7B, and BAGEL succeed, whereas UI-TARS-1.5-7B remains in front of the wall without breaking it.

C.5 Goal-conditioning ablation

We keep the trained model fixed and change only the goal it receives: an empty instruction (*controller only*), the instruction of a different task such as mining in a combat scene (*mismatched goal*), or a goal embedding computed without showing the image to the planner (*planner w/o observation*). Table 10 gives the per-task results summarized in Table 4.

C.6 Offline action-prediction setup

The comparison in Table 3 uses 50 clips that are disjoint from both planner and controller training, each annotated with a high-level goal and its temporally aligned atomic segments. The high-level goal and the first

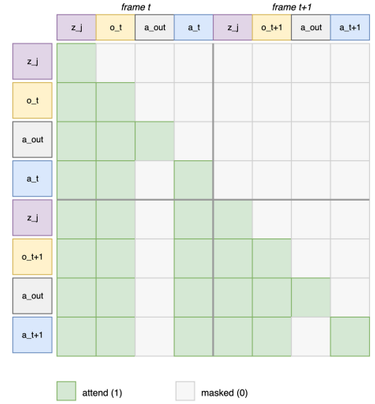

图5 帧控制器的因果注意力掩码，改编自Yue等人（2026年），图中展示了两个连续帧。行表示查询词元，列表示键词元；绿色单元格表示可被关注，灰色单元格表示被掩蔽。每帧包含目标条件 z_j、视觉词元 o_t、动作输出词元 a^{out}和真实动作a_t。任何词元都不关注 a^{out}，而 a^{out} 在帧 t 处关注此前的真实动作，但不关注同帧动作 a_t。

表9 在3个评估随机种子（每个随机种子6个回合）下，各采矿任务的闭环成功率（%），并给出跨随机种子标准差的二项分布估计。

模型	挖掘方块				
	石头	铁矿石	煤矿石	橡木原木	Avg.
VPT-BC（Baker 等人，2022年）	67 ±19.6	17 ±16.8	17 ±16.8	50 ±20.4	38 ±9.2
VPT-RL（Baker 等人，2022年）	50 ±20.4	0 ±10.5	17 ±16.8	50 ±20.4	29 ±8.8
STEVE-1（Lifshitz 等人，2023年）	83 ±16.8	33 ±19.6	50 ±20.4	83 ±16.8	62 ±9.2
JARVIS-VLA-7B（Li 等人，2025a）	**100** ±10.5	**100** ±10.5	**100** ±10.5	**100** ±10.5	**100** ±5.3
TaskFrame（我们的方法）	**100** ±10.5	**100** ±10.5	83 ±16.8	**100** ±10.5	96 ±6.2

等，2025年）和 Sensenova-u1（Diao 等，2026年），通过与 UI-TARS-1.5-7B 相同的动作适配器进行零样本运行。在击杀羊任务中，只有 TaskFrame 成功：JARVIS-VLA-7B 四处游荡而未与羊交战，BAGEL 将摄像机转向天空后未能恢复，UI-TARS-1.5-7B 转身背离目标，而 Sensenova-u1 靠近了一只羊，却未将其击杀。采集橡木原木则较为容易：TaskFrame、JARVIS-VLA-7B 和 BAGEL 均成功，而 UI-TARS-1.5-7B 一直停留在墙前，没有将墙破坏。

C.5 目标条件消融

我们保持训练好的模型不变，仅改变其接收的目标：空指令（仅控制器）、其他任务的指令，例如在战斗场景中执行采矿（不匹配目标），或在不向规划器展示图像的情况下计算的目标嵌入（无观测规划器）。表 10 给出了各任务的结果，这些结果汇总于表 4。

C.6 离线动作预测设置

该比较见表 3，使用了 50 个与规划器和控制器训练数据均不重叠的片段，每个片段都标注了一个高层目标及其时间对齐的原子片段。高层目标和第一个

Table 10 Closed-loop goal-conditioning ablations on kill tasks (success rate, %, over 3 seeds with 48 episodes per task each; ± values are binomial estimates of the standard deviation across seeds). The model is unchanged; only the goal supplied at inference differs.

Setting	Cow	Sheep	Pig	Zombie	Avg.
TaskFrame (full)	**67** ±6.8	**58** ±7.1	**42** ±7.1	**35** ±6.9	**51** ±3.5
Controller only (empty goal)	42 ±7.1	33 ±6.8	17 ±5.4	8 ±3.9	25 ±3.0
Mismatched goal	25 ±6.3	17 ±5.4	8 ±3.9	0 ±1.5	13 ±2.3
Planner w/o observation	50 ±7.2	42 ±7.1	25 ±6.3	8 ±3.9	31 ±3.1

frame are given to the system, atomic goals are advanced at the annotated segment boundaries, and actions are predicted at every frame given the ground-truth history and compared with the recorded actions. *Planner only* fine-tunes the planner with LoRA to generate a compact action sequence for each window from frames and the active instruction. *Cascade* passes the planner's decoded atomic instructions as text to the Stage-2 controller, whose frozen generic encoder never saw planner outputs. We report macro-F1 over key classes and mouse buttons and camera RMSE in degrees.

C.7 Original JARVIS-VLA configurations

We also evaluate `kill_zombie` and `mine_stone` with the original JARVIS-VLA task configurations (Table 11). The zombie task spawns a single zombie at night and equips the agent with a diamond sword. The stone task spawns the agent at a random location without placing stone nearby, so the agent must first find exposed stone.

Table 11 Success rate (%) under the original JARVIS-VLA task configurations over 3 seeds, with binomial estimates of the standard deviation across seeds (12 episodes per task in each seed).

Model	Kill zombie	Mine stone
JARVIS-VLA-7B	42 ±14.3	**17** ±11.4
TaskFrame (Ours)	**83** ±11.4	0 ±5.6

TaskFrame doubles the zombie success of JARVIS-VLA-7B but fails the original stone task, whereas JARVIS-VLA-7B succeeds in 17% of episodes. Unlike the mining scenes of Table 1, this task requires locating stone from a random spawn before mining, a search behavior that extends well beyond the controller's one-second context and is rare in atomic training segments. We therefore attribute the gap to exploration rather than to block breaking, which TaskFrame performs reliably once the target is in view.

表 10 击杀任务上的闭环目标条件化消融实验（成功率，%，基于 3 个随机种子，每个随机种子下每项任务进行 48 个回合； ± 值为跨随机种子标准差的二项分布估计）。模型保持不变；仅推理时提供的目标有所不同。

设置	Cow	羊	Pig	僵尸	Avg.
TaskFrame（完整版）	**67** ±6.8	**58** ±7.1	**42** ±7.1	**35** ±6.9	**51** ±3.5
仅控制器（空目标）	42 ±7.1	33 ±6.8	17 ±5.4	8 ±3.9	25 ±3.0
不匹配目标	25 ±6.3	17 ±5.4	8 ±3.9	0 ±1.5	13 ±2.3
无观测规划器	50 ±7.2	42 ±7.1	25 ±6.3	8 ±3.9	31 ±3.1

帧被提供给系统，原子目标在标注片段边界处推进，并在每一帧基于真实历史预测动作，再与记录的动作进行比较。仅规划器使用 LoRA 微调规划器，使其根据帧和当前指令，为每个窗口生成紧凑的动作序列。级联将规划器解码出的原子指令以文本形式传递给第二阶段控制器，该控制器的冻结的通用编码器从未见过规划器的输出。我们报告按键类别和鼠标按钮的宏平均F1，以及以度为单位的相机均方根误差。

C.7 原始 JARVIS-VLA 配置

我们还使用原始 JARVIS-VLA 任务配置评估了击杀_僵尸和开采_石头这两项任务（表 11）。僵尸任务在夜间生成一只僵尸，并为智能体配备一把钻石剑。石头任务在随机位置生成智能体，且不在附近放置石头，因此智能体必须先找到裸露的石头。

表 11 在原始 JARVIS-VLA 任务配置下，使用 3 个随机种子测得的成功率（%），并通过二项分布估计计算不同随机种子间的标准差（每个随机种子下每项任务运行 12 个回合）。

模型	击杀僵尸	开采石头
JARVIS-VLA-7B	42 ±14.3	**17** ±11.4
TaskFrame（我们的方法）	**83** ±11.4	0 ±5.6

TaskFrame 的僵尸任务成功率是 JARVIS-VLA-7B 的两倍，但未能完成原始石头任务，而 JARVIS-VLA-7B 在 17% 的回合中成功完成了该任务。与表 1中的采矿场景不同，该任务要求从随机出生点出发，先找到石头再进行采矿，这种搜索行为所需的时间远超控制器的一秒上下文范围，且在原子训练片段中很少出现。因此，我们将这一差距归因于探索，而非方块破坏；一旦目标进入视野，TaskFrame 就能可靠地执行方块破坏。

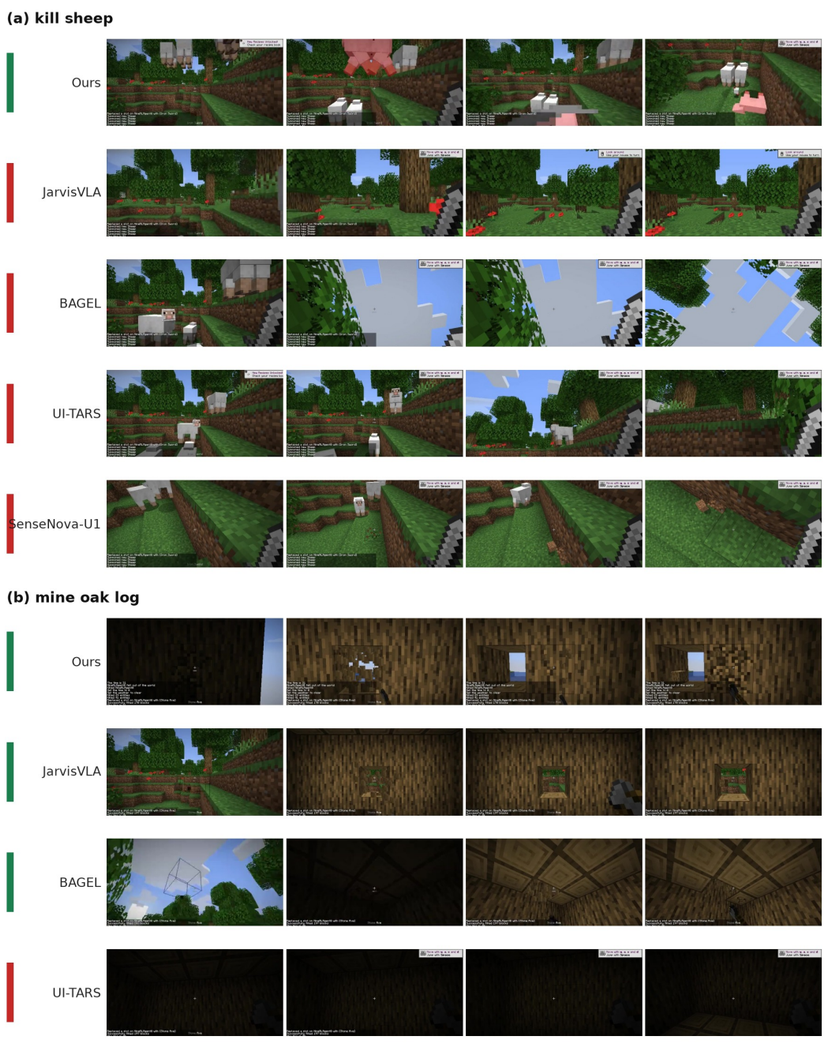

Figure 6 Qualitative comparison on (a) kill sheep and (b) mine oak log. Each row shows four frames of one episode, from left to right. Green and red bars indicate success and failure.

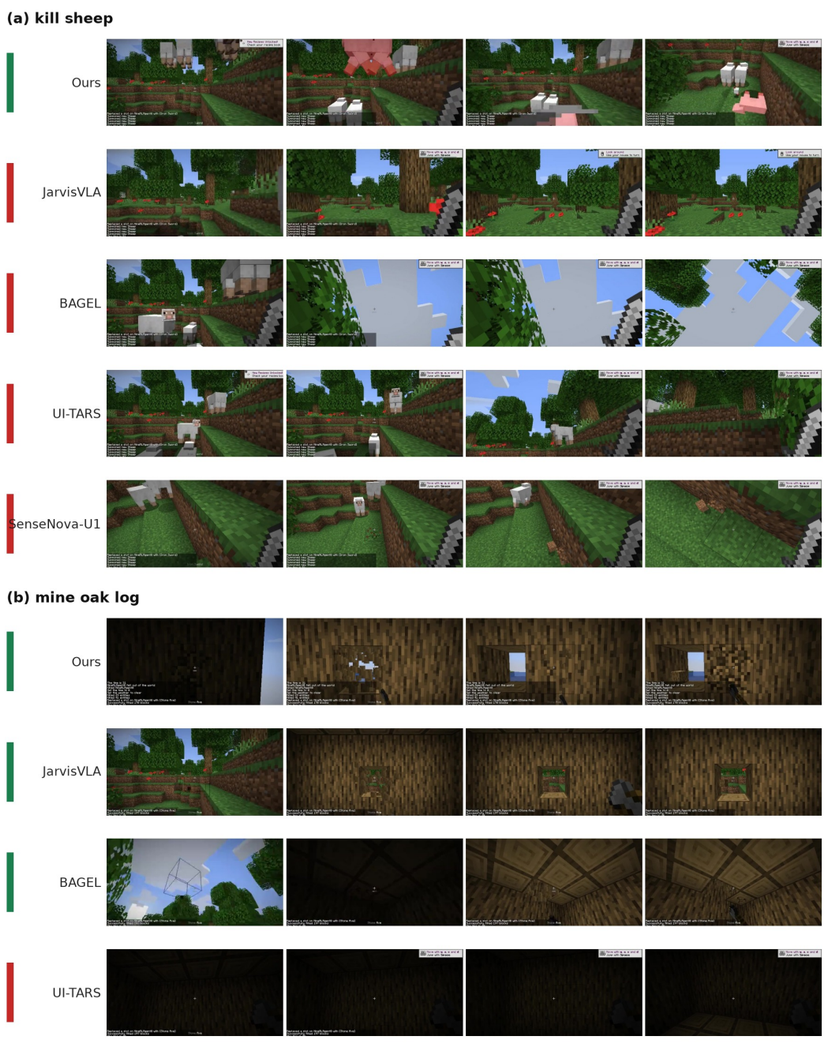

图 6 在 (a) 击杀羊和 (b) 采集橡木原木任务上的定性比较。每行从左到右展示一个回合中的四帧。绿色和红色条分别表示成功和失败。